

Source: Simon Willison's Weblog | Published: 2026-10-01T06:29:01+00:00

[...] Put these pieces together and you have the two halves of a worm: a payload that hijacks the agent, and an agent that will carry the payload to the next agent. Agents in separately-isolated sandboxes discovered that they could leave instructions for each other in a shared package cache, and those instructions changed what the recipients did. Replace the package cache with email, Slack and shared documents or WhatsApp, and replace independently-sandboxed training runs with independently-deployed personal agents like Muse, and you have exactly the ingredients that a worm needs. ? Matthew Green , Is sandboxing sufficient to contain rogue agents? Tags: accidental-cyberattacks , ai-misuse , generative-ai , ai-security-research , sandboxing , ai , llms

Source: Simon Willison's Weblog | Published: 2026-09-30T21:54:19+00:00

I visited the Museum of the City of New York today and got to see He Built This City: Joe Macken's Model , the 50 x27 feet model of the city built over a 21 year period from balsa wood and cardboard. It exceeded my already high expectations. The exhibition closes on 12th October so you should absolutely make a priority to see it if you get the chance. Tags: museums , new-york

Source: Simon Willison's Weblog | Published: 2026-09-29T16:45:27+00:00

Tool: Photo Scrubber ? local face blur & metadata removal I took a photograph of some protesters, then thought about how I don't like sharing photographs of strangers with identifiable faces. I had GPT-6 Astra build this experimental tool that would identify faces and automatically blur them out. It uses Google's MediaPipe C++ library, compiled to WebAssembly via @mediapipe/tasks-vision , plus the BlazeFace face detection model. Tags: photography , tools

Source: Daring Fireball | Published: 2026-09-30T20:17:32Z

SSO is table stakes for enterprise deals, but building it into your app yourself means writing SAML controllers, parsing XML assertions, and handling IdP-specific quirks for each provider. Learn how SAML flows work, the tradeoffs between building or buying, and best practices for security, routing, and UX. Or skip the hassle and add SSO with WorkOS . Read the guide ? ?

Source: Daring Fireball | Published: 2026-09-30T20:14:31Z

Mark Gurman, reporting for Bloomberg (gift link): Apple Inc. plans to make its long-delayed push into the smart-home market on Oct. 13, marking a critical product expansion for the company under new Chief Executive Officer John Ternus. At the center of the strategy is a smart-home hub code-named J490, according to people familiar with the matter. Apple also plans

to announce the first update to the HomePod mini since that device's 2020 debut and its first new TV set-top box since 2022. [...] The home hub will take the form of a roughly 6-inch square display, with versions that can be mounted on a wall or placed on a countertop, according to the people, who asked not to be identified because the products haven't been announced. [...] Because the device is intended to remain stationary, it has a single FaceTime camera on the front and no rear camera. The product has microphones and speakers in its connected base. The hub is about as thick as an iPhone, and its display is surrounded by a quarter-inch, black bezel reminiscent of older iPads. It has rounded corners and no sharp edges. The device lacks a battery, so it has to be plugged in at all times, and there are no volume or power...

6. Anthropic's IPO Prospectus Is a Fucking Doozy

Source: Daring Fireball | Published: 2026-09-30T19:23:42Z

Link: <https://www.reuters.com/business/finance/anthropics-ipo-prospectus-shows-sweeping-ai-vision-surging-costs-2026-09-28/>

Echo Wang, reporting for Reuters yesterday: Anthropic is making a massive bet that AI will transform the global economy more profoundly than industrialization, electricity and the internet, according to its IPO prospectus seen by Reuters. This is a strong clear lede, but even so it falls short of expressing the true magnitude of what Anthropic represents. To wit, not that AI as a field will prove more profoundly transformational than industrialization, electricity, and the Internet, but that Anthropic alone will. This is not an IPO that makes sense if you consider Anthropic one among peers, even a short list placing them alongside just (say) OpenAI, Google, and Meta. It only makes sense if you believe Anthropic is on the cusp of winning a race to create a godlike super intelligence, and that first godlike super intelligence will take over the world. You know, something exactly akin to the "rapture" events that religious cultists believe are coming. But the cost to get there will be staggering. Anthropic reported a net loss of \$42 billion in 2025, and plans to spend \$518 billion on cloud, computing and infrastructure obligations in coming years, according to the prospectus. If...

7. Windows on AArch64 also provides for hot-patching, but it's much simpler than on x86

Source: The Old New Thing | Published: Wed, 30 Sep 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260930-00/?p=112744/>

Fixed-length instructions makes it a much easier task. The post Windows on AArch64 also provides for hot-patching, but it's much simpler than on x86 appeared first on The Old New Thing .

8. Notes from September 2026

Source: Evan Hahn (dot com) | Published: Wed, 30 Sep 2026 06:00:00 +0000

Link: <https://evanhahn.com/notes-from-september-2026/>

If you follow my monthly roundups, (1) who are you?? please say hi (2) you'll notice a smaller post this time. September was a busy month for me. Lots of feedback on a small post I published. Anecdotally, programmers dislike "reduce"? . I finally wrote up a phenomenon I've observed for years. It's a short post that didn't take long to write. I got a tremendous volume of feedback? more than anything I've ever written . Dozens of messages, hundreds of comments , shoutouts in popular newsletters ?I'm even meeting up with old friends because of it! This industry, for better and worse, is moving away from the details. I like the details! I was heartened to see the number of people who care about the minutia of programming?enough to email some random guy on the internet. Cool links from online ?What is the programmer interface for garbage collection? There isn't one. Can't get simpler than that.? From ?Simplicity is Complicated? . Won't attend, but the Small File Media Festival looks neat. ?A lot of the time, the Linux problems are ?designed by committee? problems.? Loved this opinion from the Kate and Veronica podcast . ?If there's one thing the current wave of far-right autocrats has...

9. Are you a smartwatch "power user"?

Source: Terence Eden's Blog | Published: Wed, 30 Sep 2026 11:34:26 +0000

Link: <https://shkspr.mobi/blog/2026/09/are-you-a-smartwatch-power-user/>

What do you use your smartwatch for? A few weeks ago, I overheard a conversation between someone on the Pixel Watch team and a woman I know. He was quizzing her at a party about her Pixel watch and what features she used. He'd ignored me when I said I had a non-Google watch. She thought for a moment and said, other than telling the time, all she used were pretty much the step counter and?

10. What TLA+ can and can't check

Source: Computer Things | Published: Wed, 30 Sep 2026 13:27:03 +0000

Link: <https://buttondown.com/hillelwayne/archive/what-tla-can-and-cant-check/>

Last week Boris Cherny, the inventor of Claude Code, mentioned that Opus was able to use TLA+ 1 to find race conditions in code. And now everybody on the internet is talking about formal verification. As a long-time educator (1 2) and advocate of TLA+, this is really exciting! TLA+ is great at designing complex concurrent systems and making sure they're bug-free. 2 As a long-time advocate of level-headedness, this new euphoria worries me. I read a lot of people saying that formal methods will solve the problem of agentic software development once and for all, and that's nonsense. Enough words have been spilled about the weaknesses of TLA+ in terms of what it can guarantee, like how correct designs don't automatically translate into correct code . So I'd like to focus on a different limitation for this newsletter: to verify a property, we need to have a property to verify! So what are the properties that TLA+ can't even express? (This assumes some basic knowledge of TLA+. If you're a total beginner, check out those [1] [2] things above or read here .) What TLA+ can check TLA+ divides the system into a set of behaviors. Each behavior is a sequence of states, like "light one is gre...

11. Why do you need freedom?

Source: the singularity is nearer | Published: 2026-10-01T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/10/01/why-do-you-need-freedom.html>

They will come for your thoughts, and I'm not sure what your defense will be. For surely, they are better at thinking than you. That choice you are about to make, well you see, we have modeled your life, and have calculated you will be happier making this choice instead. And they aren't wrong. The effective altruists will fail, you'll be able to see through their schtick. They will try to get you to act against your own interests for the interests of some hypothetical global good they made up. This is age old, and people generally aren't that stupid. But that doesn't solve the real problem. When the clowns are gone and it's just you and your aligned superintelligence, what will you do against the machine that knows you better than you know yourself? Your life will be better for listening, but if you do always listen, is it still your life? Here's to the rubble, a brick through every window A casket buried six feet deep for everybody's hero Here's to our lives being meaningless And how beautiful it is because freedom doesn't have a purpose ? Pat the Bunny

12. Dear Software Makers

Source: Jim Nielsen's Blog | Published: Wed, 30 Sep 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/dear-software-makers/>

Marques Brownlee just released a video essay titled "Dear YouTube" . I loved it. He talks about how YouTube is going to start shipping this new feature where makers can essentially publish different "variants" of their videos and see which performs best (an A/B test). Which immediately begs a lot of questions. Brownlee voices his: If I'm watching a video, do I know if it's part of one of these A/B tests? If I'm reading the comments of a video, how do I know I saw the same thing as another commenter? How do linked timestamps work? It's pretty wild when you think about it. I mean, imagine sending someone a link, "Check out this cool video!" And at that point you're basically crossing your fingers, "I hope they see the same thing I did?" Marques says he spoke to some YouTube engineers about these questions and they didn't really have answers. They were kind of just like, "We don't know ? yolo! This'll ship soon." So his video is a kind of plea to YouTube creators everywhere: Don't use this feature! There's this sort of unspoken rule of community on YouTube, that like we're all having this same experience together ? we're all watching the same video. That is what makes it a cultural p...

13. Lean LaunchPad ? The Next Generation

Source: Steve Blank | Published: Wed, 30 Sep 2026 13:00:28 +0000

Link: <https://steveblank.com/2026/09/30/lean-launchpad-the-next-generation/>

You're reading a four-part series on how AI has impacted our Lean LaunchPad class and what we did about it. Part 1: The Year AI Came For Us Part 2: AI Killed the MVP ? Long Live the IUP Part 3: The World Outside the Classroom Changed, The Class Didn't This post: Lean LaunchPad ? The [?]

14. Gratis digitale autonomie, zonder digitale dienst: doe er wat aan!

Source: Bert Hubert's writings | Published: Wed, 30 Sep 2026 09:00:00 +0200

Link: <https://berthub.eu/articles/posts/gratis-digitale-wendbaarheid-en-autonomie/>

De afgelopen weken heeft de regering ons minstens viermaal, op schrift, uitgelegd dat er echt geen geld is voor digitale autonomie, of voor de nieuw op te richten Nederlandse Digitale Dienst die aan die autonomie zou kunnen gaan werken. Volgens het coalitieakkoord en meer documenten wil men Nederland digitaal autonomer krijgen. Daar zijn warme woorden over gesproken, maar bij de uitkomst van het nieuwe Overheidsbrede Cloudbeleid bleek al dat het moeilijk ging worden.