

1. Quoting Felix Rieseberg

Source: Simon Willison's Weblog | Published: 2026-10-05T23:56:47+00:00

Link: <https://simonwillison.net/2026/Oct/5/felix-rieseberg/>

The "old" version of Cowork runs model inference in the cloud, executing tool calls in an Anthropic-provided VM we shipped to your computer. We added the VM for capability, safety, and security reasons - mapping in just the data you explicitly added to your session. People loved what they were able to do with Claude but didn't love the disk, battery, and performance cost of running the VM locally. Also, people didn't love that closing your laptop means the work stops. The "new" version of Cowork runs model inference and the VM in the cloud. Each session gets its own sandbox, not sharing state with other sessions. When the VM needs something on the users' device (like a file), the desktop app is responsible for that file access tool call.

[...] We think this solves a lot of problems we've heard about (like using Cowork from a phone, keeping work running, or getting all the same power without losing battery to the VM) ? Felix Rieseberg , Anthropic, see also this help page Tags: claude-cowork , anthropic , claude , generative-ai , ai , general-agents , llms

2. Qwen3.8 27B addition in words

Source: Simon Willison's Weblog | Published: 2026-10-04T23:34:00+00:00

Link: <https://simonwillison.net/2026/Oct/4/qwen38-addition-in-words/>

Research: Qwen3.8 27B addition in words Colin Frasier posted on Bluesky about an experiment he ran over two years ago using GPT-4o to see how well it could "compute the sum but return the answer in words" across increasingly large numbers. Here's the chart he shared of those results: I'm confident GPT-4o didn't cheat and use a calculator, especially since it got so many of the calculations wrong, but I was inspired to run the experiment again on local hardware (a DGX Spark) to explore the effect in a fully controlled environment. I pasted his image into a Codex Remote session (GPT-6 Astra) and had it run the same experiment using Qwen3.8-27B-Q4_K_M.gguf . Here's the result for a run of 30 attempts per combination with reasoning disabled: Then I ran it again with reasoning enabled. This took a lot longer per pair, so instead of running 30 samples per square I ran just one - which results in a much less visually appealing heatmap since each square is either 100% or 0%: It got the right answer in 167 out of 169 attempts, and since these were one-shot I'm confident a second run would produce different results here. Here's a version of the report that includes the reasoning traces from...

3. pwasm 0.2a0

Source: Simon Willison's Weblog | Published: 2026-10-01T17:10:31+00:00

Link: <https://simonwillison.net/2026/Oct/1/pwasm/>

Release: pwasm 0.2a0 pwasm is one of my folly projects - an entirely vibe-coded pure Python WebAssembly engine that I built in January during my first bout of AI mania . I hadn't touched it since January, so I decided to let Claude Opus 5.5 loose on it and see if it could make any significant improvements: Evaluate current state of pwasm - then consider what it would take to get the MicroPython and micro JavaScript experiments from the research repo working under it - and what it would take to speed it up 42 commits later (with minimal follow-up prompting) it now handles almost all of the WASM specification, and the wheel from PyPI bundles working WASM builds of MicroPython , QuickJS and Micro QuickJS . I wouldn't trust this thing at all - hence the alpha version tag - but it's interesting seeing how today's models can improve on the work of models from 10 months ago. Tags: python , webassembly , vibe-coding

4. Steve Jobs, Walking Through a Mockup for Apple Park in 2010

Source: Daring Fireball | Published: 2026-10-06T02:49:44Z

Link: <https://book.stevejobsarchive.com/#photo-37>

What an apt photo, today of all days. What a gift Make Something Wonderful is. ?

5. [Sponsor] Sunnny

Source: Daring Fireball | Published: 2026-10-05T23:49:05Z

Link: <https://sunnny.com/>

How can B2B software be so friggin' delicious? Early access drops 10/19. Spread the love. ?

6. OpenAI Announces Their Text Watermarking Plans

Source: Daring Fireball | Published: 2026-10-05T22:59:14Z

Link: <https://openai.com/index/eu-text-provenance/>

OpenAI, today: The EU AI Act requires generative AI providers to make generated text identifiable in a machine-readable way. Text watermarking and detection remain early technologies with significant limitations, and views about their benefits and responsible uses are still developing. Our phased approach reflects both the EU AI Act requirements as well as the technology's limitations, with an emphasis on transparency about what a text watermark can and cannot tell people: Starting today, API customers globally will be able to opt in to text watermarking for select models. Text watermarking will remain off by default in the API. Over the coming weeks, we will add an invisible watermark to eligible ChatGPT and Codex text output in the European Union. We're opening applications to access our text watermark detector. Access will initially be limited to approved researchers and expert organizations that can help us evaluate and improve the technology. First, unlike Anthropic, OpenAI is only forcing this upon users in the EU, where it's mandated by their ill-considered 2024 regulation. Second, also unlike Anthropic, OpenAI has made this available via an API so developers should be able...

7. If somebody tries to hot-patch an already-hot-patched function, how do they avoid conflicts?

Source: The Old New Thing | Published: Mon, 05 Oct 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20261005-00/?p=112755/>

Hot-patching is for Windows, not for you. The post If somebody tries to hot-patch an already-hot-patched function, how do they avoid conflicts? appeared first on The Old New Thing

8. Benchmark In Milliseconds

Source: matklad | Published: 2026-10-05T00:00:00+00:00

Link: <https://matklad.github.io/2026/10/05/benchmark-milliseconds.html>

How long should a micro benchmark run? My rule of thumb is to tweak the input size until the benchmark takes about 300ms, for the following reasons:

9. The indicators of a quality website (somewhat silly)

Source: Maurycy's Blog | Published: Sun, 04 Oct 2026 00:00:00 +0000

Link: <https://maurycyz.com/misc/quality/indicators.html>

You know it's a good website when it: Uses subdomains (except for username.platform.com) ... especially multiple levels Has a non-standard port number in URL Is served over a plaintext connection Uses very nearly default styles All tag names are written in SCREAMING CASE The path ends in .html The URL has meaningful directories and names The HTML has typos Has commented out draft work Handmade index pages --> Of course, these are all just indicators of age. 1. The use of subdomains fell out of favor due to IP address exhaustion and reverse proxy software. If your going to be setting up a proxy anyways, it doesn't make much sense to fiddle around with DNS. (which often makes setting up TLS harder) On my site, the server handling / and /mcu aren't even on the same continent. At one point, it was standard practice to setup a subdomains per protocol in case one had to be moved to new infrastructure, but HTTP is the last one standing: public access FTP/Gopher/Telnet are just about extinct. Even in places where they absolutely should be used, subdomains seem to have fallen out of fashion. Many companies will register a new domain for an internal service instead of putting it under their...

10. [RSS Club] Changes to the RSS and Atom feeds

Source: Terence Eden's Blog | Published: Mon, 05 Oct 2026 11:34:03 +0000

Link: <https://shkspr.mobi/blog/2026/10/rss-club-changes-to-the-rss-and-atom-feeds/>

This content is exclusive to RSS and Atom subscribers. It is not on the web. Shhhhhh! A few weeks ago I accidentally broke some of your feed readers. I was using blocks in the feed. That allows me to put in raw HTML without the need for escaping. Unfortunately, some feed readers unescaped text which should have remained escaped. Confusing, innit? Anyway, if you've a mind to, I'd b?

11. Missing Link

Source: Tedium: The Dull Side of the Internet. | Published: 2026-10-06T02:59:08Z

Link: <https://feed.tedium.co/link/15204/17490109/link-newspaper-virginian-pilot-remembrance>

Some spare thoughts about Link, a newspaper that turned 20 years old today?but more or less changed my life. It's been 18 years, but I still miss that job.

12. Pluralistic: Scrutinized (05 Oct 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 05 Oct 2026 18:58:11 +0000

Link: <https://pluralistic.net/2026/10/05/pervert-glasses/>

Today's links Scrutinized: Pervert glasses vs the authentic self. Hey look at this: Delights to delectate. Object permanence: DRM-free in the UK; Maple-leaf roses; Italian Wikipedia shuts down; Vancouver's property speculator tax works; Ex-Soviet leftists launch "September"; "The Every"; Nielsen x Chinese internet; Get-a-cluevog; Cat EULA; Heart carafe; Millennials screwed; Apple's unlawful evil. Upcoming appearances: Virtual, South Bend, Hudson, Calgary, Winnipeg, Paris, OVancouver, Victoria, Ottawa, Kilkenny, Oxford, Montreal. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Scrutinized (permalink) We're repeatedly told that the killer app for pervert glasses ? Meta's AI-enabled camera glasses and their imitators ? is that they will remind you of someone's name and spare you the social awkwardness of not recognizing them. I admit to finding this proposition very seductive. I'm mildly face-blind (I had two aunts I wasn't able to reliably tell apart until I was 12 or 13 and one of them changed her hair), and I meet a lot of people and often find myself incredib...

13. We are going to kill ?unalive?

Source: Anil Dash | Published: 2026-10-06T00:00:00Z

Link: <https://anildash.com/2026/10/06/kill-unalive/>

I talk to a lot of old people, those who were born in the 20th century, and if I ask them what the word ?unalive? means, they usually have no idea what I'm talking about, except for some of them who have kids or who study Internet culture. This will, of course, probably seem very weird to most people who were born in or grew up in the 21st century. Just to recap for the olds: ?unalive? is the word you use to represent concepts like dying, or death, or killing or being killed, on digital platforms where saying those words accurately will cause the algorithm to punish or censor you. Or, maybe, where the perception is that using those words will result in being censored by the algorithm, and no one is actually willing to find out what happens if you use the forbidden words. This sort of attack on people's expression started on platforms like TikTok, where nearly all content is distributed through an algorithmic feed, but has since become ubiquitous in nearly all digital media that we see. In fact, these tics are now so prevalent that it's routine to hear people using this kind of language in everyday life, even though there's not yet an algorithm to appease in the physical world. I?v...

14. How fast is Python 3.15?

Source: Miguel Grinberg's Blog | Published: Mon, 05 Oct 2026 10:13:17 GMT

Link: <https://blog.miguelgrinberg.com/post/how-fast-is-python-3-15>

It's October once again, and that means it is time to take the new release of Python for a spin (technically, it is the 3.15.0rc3 release that I'm using, the official 3.15 release is still a few days out). As I did with my Python 3.14 performance article of a year ago, today I'm

sharing a new run of my informal Python benchmark, comparing Python 3.15 against previous interpreters all the way back to 3.10. If you are not interested in the charts and the tables and just want to read my analysis, feel free to jump to the conclusions section at the end.