RSS Digest
Generated at: 2026-10-12 15:02:21

1. llm-openai-decisions 0.1a0
Source: Simon Willison’s Weblog | Published: 2026-10-06T23:04:13+00:00
Link: https://simonwillison.net/2026/Oct/6/llm-openai-decisions/
Release: llm-openai-decisions 0.1a0 OpenAI released their new Jev-style Decisions API , as previously announced at last week’s DevDay. Since I already have an llm-typesafe plugin for talking to Jev, I had GPT-6 Astra read the new OpenAI API documentation and build an llm-openai-decisions plugin inspired by llm-typesafe . Unlike Jev, the new gpt-6-luna decision model supports image input in addition to text. Both models charge for input it and not for output: OpenAI’s is 10 cents per million input tokens, Jev’s is 4.2 cents per million. Otherwise the API shape is very similar to Jev, at least conceptually. Jev supports three question types for yes/no, choices, or scores. OpenAI Decisions supports the same three types. Install the plugin like this: llm install llm-openai-decisions Here’s an example query against an image attachment: llm -m openai-decisions/gpt-6-luna \ -a https://static.simonwillison.net/static/2025/two-pelicans.jpg \ -s ’ Does this image contain any mammals? ’ And example output: { "type" : " predicate " , "name" : " evaluation " , "probability" : 0.0 } Consult the README for full details of how to run the other types of questions. Tags: openai , llm , coding-agent...

2. llm-mistral 0.16
Source: Simon Willison’s Weblog | Published: 2026-10-06T21:32:34+00:00
Link: https://simonwillison.net/2026/Oct/6/llm-mistral/
Release: llm-mistral 0.16 Adds support for reasoning models, such as the newly released Mistral Large 4 . Tags: llm , mistral , llm-reasoning

3. EmbeddingGemma 2
Source: Simon Willison’s Weblog | Published: 2026-10-06T20:37:53+00:00
Link: https://simonwillison.net/2026/Oct/6/hn-49983751/
My comment on EmbeddingGemma 2 ? Hacker News. I really appreciate that EmbeddingGemma 2 is under the Apache 2.0 license. For embedding models in particular, I don’t think it makes sense to use a closed, proprietary, hosted-only model. Most applications of embedding models involve calculating thousands or even millions of embedding vectors and storing them for later comparison. If your model is proprietary, the vendor is likely someday going to decide to stop offering that model. They’ll have a better model to replace it, but you still need to pay to re-calculate those millions of stored existing vectors. (In April 2024 OpenAI offered to "cover the financial cost of users re-embedding content with these new models" - https://openai.com/index/gpt-4-api-general-availability/ - but I don’t think that’s something we can rely on from every provider.) Notably, I don’t want to host the model myself . I’d much rather pay a provider for a hosted model while knowing that if they ever stop hosting it I can run the open weights version myself - or find another vendor who can do that for me. Tags: google , ai , generative-ai , embeddings , gemma

4. ? X Once Again Marks the Spot
Source: Daring Fireball | Published: 2026-10-11T21:05:33Z
Link: https://daringfireball.net/2026/10/x_once_again_marks_the_spot
If you returned from exile on a desert island and took a look at X today for first time in a year, you might reasonably guess that Elon Musk got bored and sold it while you were away.

5. Apple Only Started Suggesting iOS 26 Users Upgrade to 27 Three Days Ago
Source: Daring Fireball | Published: 2026-10-11T19:02:38Z
Link: https://www.macrumors.com/2026/10/08/apple-now-pushing-ios-27/
Joe Rossignol, on Thursday: Starting today, iPhone users still running iOS 26 will see iOS 27.0.1 as the prominent software update in the Settings app under General ? Software Update,

with iOS 26.7.1 now relegated to a smaller ?Also Available? section at the bottom. Just the other day I posted about how most users only get OS updates when Apple pushes them through automatic updates. But even for users who update manually, until Thursday, the default suggested update for users on iOS 26.x was to remain on the latest version of iOS 26. To get iOS 27 you had to scroll down ?below the fold? and most users don?t do that. iOS 27 adoption amongst normals is very slow, but that?s only because Apple wants it slow. That?s not because iOS 27 is buggy ? it speaks to its Snow-Leopard-like focus on bug fixes and performance updates that it?s actually remarkably stable. The obvious explanation is that Apple is slow-rolling 27 updates to gate access to Siri AI. ?

6. Why I Get Bored
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Sun, 11 Oct 2026 16:00:00 GMT
Link: https://idiallo.com/byte-size/why-i-get-bored
For me, boredom is just a signal. It means I've hit a roadblock and my brain is trying to convince me not to go any further. "Give up. This sucks anyway. Why are you still wasting your time on this silly thing?" That's how I feel when I do things I'm really bad at. I think I'm not good at making YouTube videos. Whenever I think of a subject I find really exciting, I start writing, record a first pass, and then load DaVinci Resolve. With my current schedule, job, family, and responsibilities, I've found that it takes me a whole week to polish three minutes of video. That means I hear the same material I wrote and recorded over and over until it becomes completely boring. After a few weeks, I lose interest in the subject and have no motivation to finish it. But I think that's just an excuse I tell myself. I like the videos I've managed to complete. I wish I could change some parts, but I still like the final product. It's just that the process gets in the way, and it does so because I have no expertise in making videos. But then again, I always go back to my trusted method: if you're bad at something, do it at least a hundred times before you decide to give up. We all suck at the be...

7. Don't diagnose your opponents
Source: seangoedecke.com RSS feed | Published: Sun, 11 Oct 2026 00:00:00 GMT
Link: https://seangoedecke.com/dont-diagnose-your-opponents/
When you disagree with someone, there are two paths forward. The first is to try to engage with what they?re saying: what their reasons are, whether their position is true, what evidence could be brought to bear, and so on. The second is to diagnose them: to engage with the second-order question of why they?re saying it. For instance, if you?re arguing with somebody about microservices 1 , you could say ?well, you only like microservices because you want to polish your resume?, or ?you only like monoliths because you?re scared of distributed systems?. This happens a lot in online arguments. On the question ?are AI models improving?, some anti-AI people will argue that people only think so because AI psychosis is rotting their brain , and the dumber they get, the smarter AI models seem in comparison. Likewise, many pro-AI people argue that people only disagree because they?re Luddites, or scared of change, and so on. None of these arguments are about the quality of AI models. It feels so good to diagnose your opponents because there?s no possible response. If you say ?I think AI models are improving because these benchmark scores are rising?, somebody might point out that the AI la...

8. One way of thinking about access specifiers in C++/WinRT implementation types
Source: The Old New Thing | Published: Fri, 09 Oct 2026 14:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20261009-00/?p=112764/
Mapping them to concepts from C# and the now-defunct C++/CX. The post One way of thinking about access specifiers in C++/WinRT implementation types appeared first on The Old New Thing .

9. Just say no to OpenAI and Anthropic
Source: micahflee | Published: Sun, 11 Oct 2026 23:03:05 GMT
Link: https://micahflee.com/just-say-no-to-openai-and-anthropic/

I recently read End Times Fascism and the Fight for the Living World by Naomi Klein and Astra Taylor. It’s an excellent book and I highly recommend that everyone read it. Tech oligarchs make up an important component of the end-times-fascist coalition. One thing that stood

10. I paid people to try and follow my README
Source: Terence Eden?s Blog | Published: Sun, 11 Oct 2026 11:34:39 +0000
Link: https://shkspr.mobi/blog/2026/10/i-paid-people-to-try-and-follow-my-readme/
Developers suck at writing documentation. We like to say RTFM but often there ain’t no FM to R. At best we have a half-arsed collection of notes, some obsolete tips, and a link to a desolate forum full of people asking the same questions again and again. That’s no way to treat our users. I naïvely believe there’s a better way. As I’ve written before, you should sit down and actually test your ?

11. The Fascinating World of Domains
Source: Jim Nielsen?s Blog | Published: Sun, 11 Oct 2026 19:00:00 GMT
Link: https://blog.jim-nielsen.com/2026/domains-article-on-wired/
Pete Millspaugh has a lovely article over on WIRED that peers into the world of domains. (If you?re as interested in the esoteric details of domains as I am, follow Pete?s writing at Dot Com Press .) You should just read the article, but here are a few things that stood out to me. The .com Registry Did you know who the registry for .com is? I didn?t. The registry that owns .com, Verisign, makes a billion and a half dollars every year just by keeping the lights on. To Verisign?s credit, the lights never go out: It reports 100 percent uptime over the past 28 years. It?s a pretty slick business model, watching over a database of words and collecting annual renewal fees. Warren Buffett thought so too, investing in Verisign in 2012, when its stock price was about one-eighth what it is today. Why is it simultaneously surprising and not surprising that the company running .com makes billions and has 100% uptime over almost three decades? Maybe GitHub should reach out for some advice, lol. Buying Domains Is a Spy Thriller Apparently the last time you could buy generic top-level domains (gTLDs) was in 2012. So now that ICANN is providing another window of purchases in 2026 ? and nobody kno...

12. Refurb weekend II: Alpha Micro Eagle 300 and friends
Source: Old Vintage Computing Research | Published: 2026-10-11T18:47:19.668-07:00
Link: https://oldvcr.blogspot.com/2026/10/refurb-weekend-ii-alpha-micro-eagle-300.html
This Refurb Weekend is a sequel in two ways. Six years ago I fixed the power supply and case fan in my 68EC030-based Alpha Microsystems Eagle 300, a more modernized and substantially faster Alpha Micro running AMOS, their unique DEC-inspired multiuser operating system. These blue-collar boxes were very common in many vertical markets due to their tightly integrated business-specific software packages, and are especially unusual for running their big-endian 68000-series processors little -endian (for data file backwards compatibility with their WD-16 ancestors). The 1990s-era Alpha Micro Eagle family improves on the originals with later 68K and even ColdFire CPUs, plus Ethernet and faster SCSI. The second way it’s a sequel is to our previous Alpha Micro refurb weekend , a double header restoring an Alpha Micro AM-1000 and AM-1200, two members of Alpha Micro’s most popular series of computers and some of the closest things they made in the early ’80s to desktop workstations. From the outside, the Eagle comes in a PC-style case, but it’s not a PC: This particular Eagle 300, with a 40MHz 68EC030 CPU, runs ampm.floodgap.com using the webserver provided with AlphaTCP ? yes, Alpha Micros...