RSS Digest
Generated at: 2026-10-30 14:52:01

1. API credits benefit for eligible ChatGPT subscriptions
Source: Simon Willison’s Weblog | Published: 2026-10-29T22:48:29+00:00
Link: https://simonwillison.net/2026/Oct/29/api-credits/
API credits benefit for eligible ChatGPT subscriptions Earlier this month Anthropic announced a generous new benefit for their Max and Team subscriptions: $100, $200, and $500 per month subscribers would get the equivalent quantity of (expiring) credits to use with the Anthropic API. OpenAI have now launched pretty much exactly the same deal: if you have a $100, $200, or $500 subscription plan with them you now get that same quantity of credits to apply to one of your API organizations. I maintain subscriptions with both companies, so this is great news for me! I may even ship some public-facing demos now that the credits are already covered by my existing subscriptions. Tags: ai , openai , generative-ai , llms , anthropic , llm-pricing

2. Building a multiplayer game world with autonomous Fable-class models
Source: Simon Willison’s Weblog | Published: 2026-10-29T17:00:26+00:00
Link:
https://simonwillison.net/2026/Oct/29/how-i-built-the-most-ai-pilled-company-in-the-world/
How I Built the Most AI-Pilled Company in the World Don’t be put off by the clickbait headline, this piece by Jacob Sansbury describing how they are building Spawn is a fascinating exploration of the bleeding edge of what’s possible with agentic engineering against frontier models. Spawn is a massively multiplayer gaming environment where players can use prompts (to the in-game agent, Savi) to modify the game world. I fired it up on my phone for a few minutes and was able to prompt a castle into existence, then a dragon, then ride the dragon to fly around that castle. By Jacob’s telling, Spawn is now being developed autonomously by a swarm of Fable-powered agents costing over $3 million in tokens per month ("we spend 30x payroll on AI tokens"). Our agents change the live code directly and ship to production hourly, so they read every error, bug report and metric themselves. About 60% of our work starts from a production error, a Savi bug report (Savi is the AI you build games with in Spawn), something odd in our logs, or a feature request. Savi files bugs herself. When a user hits one, or Savi hits one in the product, she files it. It goes straight to an agent that decides whether...

3. Quoting Ben Hylak
Source: Simon Willison’s Weblog | Published: 2026-10-29T14:06:27+00:00
Link: https://simonwillison.net/2026/Oct/29/ben-hylak/
every product i use should just be a headless mcp, that feeds into an agent, where i never have to think about anything underneath. except for our product, of course. in our product, customers love the ui and love us very much. we’re special that way. ? Ben Hylak Tags: model-context-protocol , ai

4. You can just ask AI to do things
Source: seangoedecke.com RSS feed | Published: Thu, 29 Oct 2026 00:00:00 GMT
Link: https://seangoedecke.com/you-can-just-ask-ai-to-do-things/
One of the most important skills in working with AI is agency : being willing to try new and ambitious things. This is true for what you build, but it?s also true for how you build. AI models are improving so quickly that most of us are bottlenecked 1 on noticing that AI agents can help automate particular tasks. Here are some examples of that from my daily work. Asking the agent to run my dev server. Agents are very good at fixing weird dev environment issues because they work fast, have infinite patience, and have a huge breadth of knowledge. I almost never run the bash commands to start my dev server anymore. Asking the agent to make requests to my dev server. Part of why I don?t run my dev server manually is that I don?t make requests to it manually either. Instead I?ll say ?okay, now start the dev server and run through the happy paths with curl?. The relatively manual process of assembling a set of curl requests and validating the output can now be entirely automated 2 . Asking the agent to watch as I make a

production change. Every time I do a deploy or flip feature flags, I ask an agent to poll metrics and logs before and after the change. Agents can ingest more logs than...

5. The support case of a customer who said that all they could see was a black screen
Source: The Old New Thing | Published: Thu, 29 Oct 2026 14:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20261029-00/?p=112808/
Then again, that is what they asked for. The post The support case of a customer who said that all they could see was a black screen appeared first on The Old New Thing .

6. Apple Drops App Store Commission by a Whopping 4 Percentage Points in Taiwan
Source: Daring Fireball | Published: 2026-10-29T21:03:42Z
Link: https://developer.apple.com/news/?id=w9mxlhek
Apple Developer News: Apple is making changes to the Taiwan storefront of the App Store, following discussions with the Taiwan Fair Trade Commission. As of October 31, 2026, the commission rate for standard Apple In-App Purchase and paid app transactions will be 26% on the Taiwan storefront of the App Store on iOS and iPadOS. Currently, that rate is 30%. Well there we go. The App Store in Taiwan ? just like in Japan ? is fixed. ?

7. A Health Update From Craig Hockenberry
Source: Daring Fireball | Published: 2026-10-29T19:57:21Z
Link: https://furbo.org/2026/10/28/a-health-update/
Craig Hockenberry: If you noticed that the Iconfactory sites were kinda broken last week, now you know why. It?s because I was kinda broken last week. Now I?m finishing this post from a comfy chair on our deck. Typing into the latest version of WordPress and PHP. And I couldn?t be happier. No, you teared up a little reading this. Very good news from a very good friend. ?

8. Rolfe Winkler at The WSJ: ?The Beleaguered Chip Sellers Getting Their Revenge on Apple?
Source: Daring Fireball | Published: 2026-10-29T19:44:11Z
Link: https://www.wsj.com/tech/apple-iphone-memory-chips-ai-c0084cc9?st=FNAyaT
Rolfe Winkler, in a crackerjack story for The Wall Street Journal (gift link): Blevins, who left Apple in 2022, sometimes set up back-to-back meetings with suppliers so that rival sales teams might run into each other, stoking competitive tension. Sales representatives learned to pack nutrition bars to sustain them through negotiating sessions that Apple executives sometimes pushed past midnight without a dinner break. Blevins?s bad-cop reputation gave his staff leverage: They threatened to call him in if contract terms weren?t to their liking. One significant contract term allowed Apple to audit its suppliers? books. Flash makers typically had to designate Apple as a ?most favored customer,? which guaranteed that no other buyer got a lower price, according to people familiar with the agreements. Apple?s audits allowed it to verify that no one paid less. Apple caught Micron selling flash at cheaper prices to another customer and won an arbitration settlement worth tens of millions of dollars more than a decade ago, said people familiar with the dispute. Apple audited its suppliers so often that Blevins negotiated a volume discount with the auditor, Deloitte. The whole story is a f...

9. "No way to prevent this" say users of only language where this regularly happens
Source: Xe Iaso’s blog | Published: Thu, 29 Oct 2026 00:00:00 GMT
Link: https://xeiaso.net/shitposts/no-way-to-prevent-this/memory-safety/CVE-2026-98432/
In the hours following the release of CVE-2026-98432 for the project Linux , site reliability workers and systems administrators scrambled to desperately rebuild and patch all their systems to fix a use-after-kfree where retransmitting a crafted connection request while finishing the handshake of that same connection causes funny things to happen up to and including arbitrary code execution. This is due to the affected components being written in C, the only programming language where these vulnerabilities regularly happen. "This was a terrible tragedy, but sometimes these things just happen and there’s nothing anyone can do to stop them," said programmer Prince Julien Hoppe, echoing statements expressed by hundreds of thousands of programmers who use the only language where 90% of the world’s memory safety vulnerabilities

have occurred in the last 50 years, and whose projects are 20 times more likely to have security vulnerabilities. "It’s a shame, but what can we do? There really isn’t anything we can do to prevent memory safety vulnerabilities from happening if the programmer doesn’t want to write their code in a robust manner." At press time, users of the only programming lang...

10. Custom Voices With Qwen3-TTS
Source: Armin Ronacher’s Thoughts and Writings | Published: 2026-10-29T00:00:00+00:00
Link: https://lucumr.pocoo.org/2026/10/29/qwen3-custom-voices/
With my kids I?m currently building a little dungeron crawler and we?re trying to learn a bit about game development along the way. However given that we live in the modern times of increasing capabilities we wanted to get a narrator and voice over. But we also wanted it to not sound like everyone?s standard AI voices so I looked into what options there are on the table. After some investigation I narrowed down to using Qwen3-TTS which uses voice embeddings for custom voices. With a bit of creativity you can also do more with those custom voices than appears initially. In case this is of interest to you, this post just covers what we learned and with some examples. You can probably point your agent to it and it can do something with it. What Qwen Offers Qwen3 is not the only TTS model in town, and it wans?t even my first choice. We initially used Kokoro . However Qwen3 supports voice cloning and we wanted to have it feel a bit more unique. In practice it?s a bit tricker though. Qwen3 ships three kinds of models: 1.7B-Base (and 0.6B): clones a voice from a few seconds of reference audio 1.7B-CustomVoice : comes with nine built-in speakers and instructions 1.7B-VoiceDesign : invents...

11. Avoid Public Functions Calling Other Public Functions
Source: matklad | Published: 2026-10-29T00:00:00+00:00
Link: https://matklad.github.io/2026/10/29/pub-calls-pub.html
A small-scale rule of thumb. This article will land much better with a concrete example, but I am drawing a blank here, so you gotta just blindly trust my say so!

12. A New Kind of Chemistry
Source: Construction Physics | Published: Thu, 29 Oct 2026 15:45:10 GMT
Link: https://www.construction-physics.com/p/a-new-kind-of-chemistry
Thanks to Kyle Schiller at Marathon Fusion for conversations and information that informed this essay.

13. Celebrating SIM Cards
Source: Jim Nielsen?s Blog | Published: Thu, 29 Oct 2026 19:00:00 GMT
Link: https://blog.jim-nielsen.com/2026/celebrating-sim-cards/
It?s time to celebrate a little win. The other day, my mom contacted me saying her 10-year-old iPhone wouldn?t turn on. After getting a hold of the device and doing some family IT support, I concluded the battery was dead. So I bought her a new (old) one on Walmart (SE 3rd gen., if you must know). Her old phone used a SIM card, so I was expecting the process to go something like this: Take the SIM out of the old phone Put it in the new phone Setup the device Troubleshoot why there was no service Call my Mom?s carrier and work through the issue But step 4 and 5 never happened. The new phone ?just? connected. I called her number. I texted it. Everything worked as expected. Everything worked as expected. Everything worked as expected. I didn?t have to Google a single question. I was floored. Flabbergasted. I know I shouldn?t have been because that?s how SIM cards are supposed to work, right? Like that?s the whole promise of a SIM card: you can seamlessly swap them between phones. And in this case, the tech worked exactly as promised. As I write this, I?m kind of amazed at how amazed I am that everything ?just worked?. Shout out to anyone who has ever worked on SIM card technology. Mo...

14. What I learned about power from LBJ and Robert Moses
Source: Dwarkesh Podcast | Published: Thu, 29 Oct 2026 16:48:39 GMT
Link: https://www.dwarkesh.com/p/lessons-in-power

Preparing to interview Robert Caro.

15. Extremely brief history of PDP computers
Source: Entropic Thoughts | Published: Thu, 29 Oct 2026 00:00:00 +0100
Link: https://entropicthoughts.com/history-of-pdp-computers
I don?t know why I started to put this together. These machines fascinate me, because so much of early computing/hacker type culture happened on them. I started writing a table, but realised a timeline would be more intuitive. In this graph, the price is for a minimal configuration and adjusted for inflation. The time is the memory cycle time, a decent proxy for speed. Some notes: The PDP-1 was created mostly as an experiment. At the time, industrial control systems used either humans or complicated contraptions of logic boards to make decisions. Some engineers at DEC wanted to see if there was a customer base willing to use general-purpose computers instead. The PDP-4 was a variant of PDP-1 that was a little slower and much cheaper. It was not compatible with the PDP-1. The PDP-4 was upgraded with new chip technology and became the PDP-7, which was then optimised further for cost and speed and became the PDP-9. The PDP-5 was created when they realised the PDP-1 and PDP-4 machines were overkill for industrial control systems. By being smaller, it could be much cheaper, at the cost of compatibility with PDP-1 and PDP-4. It was still intended for roughly the same market. The chip up...

16. Making agents fast - what I’ve learnt over the last 12 months
Source: Martin Alderson | Published: Thu, 29 Oct 2026 00:00:00 GMT
Link: https://martinalderson.com/posts/making-agents-fast/?utm_source=rss&utm_medium=rss&utm_campaign=feed
Practical lessons on speeding up AI agents: inference speed (TTFT vs tok/s), tool call traces, building the tooling agents want, and hardware and sandboxing.