RSS Digest
Generated at: 2026-11-03 15:40:27

1. I discovered a sound that leaks your exact location (and it’s everywhere) - YouTube
Source: Simon Willison’s Weblog | Published: 2026-11-03T05:32:25+00:00
Link: https://simonwillison.net/2026/Nov/3/a-sound-that-leaks-your-exact-location/
I discovered a sound that leaks your exact location (and it’s everywhere) - YouTube Genius piece of OSINT by Colsto, who demonstrates how you can take almost imperceptible background noise of planes from a video soundtrack, figure out their Doppler curves, then cross reference them against ADS-B data to find the exact location of the video based on their speed, distance, and gaps between different flights. Via Lobste.rs Tags: privacy , youtube

2. October sponsors-only newsletter
Source: Simon Willison’s Weblog | Published: 2026-11-03T05:28:59+00:00
Link: https://simonwillison.net/2026/Nov/3/october-newsletter/
I sent out the October edition of my sponsors-only monthly newsletter . If you are a sponsor (or start a sponsorship now) you can access it here . This month: The October inflection point for vibe game development The vulnerability apocalypse gets hot LLM releases API credits for $100+/month subscriptions What I’m using right now My software releases My September newsletter can now be seen here , as a preview of what you’ll get. Pay $10/month to stay a month ahead of the free copy! Tags: newsletter

3. llm 0.37
Source: Simon Willison’s Weblog | Published: 2026-11-02T23:49:30+00:00
Link: https://simonwillison.net/2026/Nov/2/llm/
Release: llm 0.37 There’s a lot of stuff in this new release of my LLM CLI tool and Python library, representing around two months of accumulated features and fixes. I added gpt-6.1-sol and gpt-6.1-terra , plus a new llm openai endpoint add command to make it easier to configure additional models that use OpenAI-compatible endpoints. I switched the default model (if you have not set your own default) to gpt-6-luna - half the price of the previous default, gpt-5.6-luna . LLM’s support for embeddings now accepts that different binary content may have different content-type - necessary for models like Gemini Embedding 2 , which can embed text, audio, images, and video. A lot of open source maintainers complain about a flood of AI-assisted PRs. LLM is my project that has been most affected by this and, to be honest, I’m not finding it to be a problem. That’s mainly because I have Opus 5.5 or GPT-6 Astra help me with triage to spot the most useful fixes! This release includes fixes from 12 authors - almost all AI-assisted - and the burden of reviewing them wasn’t nearly as great as I expected. I really see them as bug reports with attached examples - sometimes I land them, sometimes I...

4. Apple Announces Record $114.5B Fourth-Quarter Results
Source: Daring Fireball | Published: 2026-11-03T00:10:53Z
Link: https://sixcolors.com/post/2026/11/apple-announces-record-114-5b-fourth-quarter-results/
Jason Snell, Six Colors: On Monday, Apple announced record fourth-quarter earnings , with total revenue of $114.5B, up 12 percent from the year-ago quarter. iPhone revenue was up 17.5%, Mac revenue was up 23.7%, and Services revenue was up 9.8%. Wearables revenue was down 7.9% and iPad revenue was down 10.4%. [...] And now, to help you visualize what Apple just announced, here is our traditional barrage of charts and graphs. As usual, Snell also has a transcript of the analyst call ? Ternus?s first as CEO. My favorite tidbit was his line at the end of this answer: Richard Kramer, Arete Research: Your latest MacBook Pros included more iOS features like Touch and Dynamic Island as well as, obviously, Siri AI. Should we anticipate further close linking of macOS and iOS features, maybe even pointing at some point to a single unified OS, either for simplicity or to reduce costs? John Ternus: Look, we are very proud of both the Mac and the iPad, and I think what you?ve seen over the years is we?ve really enjoyed the ability to bring features and capabilities across the ecosystem and create really compelling features that span things between handoff between iPhone and Mac and iPad, and s...

5. [Sponsor] OurCal
Source: Daring Fireball | Published: 2026-11-02T22:43:59Z
Link: https://ourcal.com/?utm_source=daringfireball&utm_medium=sponsorship&utm_content=rssitem
?I didn?t see you put that in the diary.? Sound familiar? Shared calendars solved scheduling, not communication. OurCal connects any calendar, new or existing, to a group chat. Changes post there automatically, so nobody misses them. Because it?s a messenger, we made everything End-to-End Encrypted, certified by France?s cybersecurity agency. It?s native SwiftUI across Apple?s platforms, and the App Store just featured it in ?Great apps on iPhone Duo?. We?re Apple people, but we think it?s the best shared calendar on Android too. Free to try, then a low monthly group subscription. No ads, no tracking. ?

6. Apple Introduces New Base MacBook Pro With M6, 14-Inch-Only
Source: Daring Fireball | Published: 2026-10-27T17:00:46Z
Link: https://www.apple.com/newsroom/2026/10/apple-introduces-the-new-14-inch-macbook-pro-with-m6/
Apple Newsroom?s ?Pro and Tell? announcement week continues: Apple today announced the new 14-inch MacBook Pro with M6 , delivering even more performance, powerful on-device AI capabilities, and faster connectivity for college students, business users, creators, and more. M6 features an advanced 12-core CPU with the world?s fastest CPU core, so everything from productivity tasks to more demanding pro workflows feels incredibly fast and responsive. The 12-core GPU delivers up to 30 percent faster graphics and up to 25 percent faster AI workloads when compared to M5, making the 14-inch MacBook Pro with M6 a compelling value for users who need serious power. For those upgrading from M1, the performance gains are even more significant, with up to 11×faster graphics and up to 13×faster AI performance. With additional Neural Accelerators in the GPU and a Dual 16-core Neural Engine, M6 turns MacBook Pro into an even more powerful platform for AI ? from taking advantage of Siri AI and Apple Intelligence features across apps and macOS to running LLMs locally, all while preserving user privacy. In addition, the 14-inch MacBook Pro features Wi-Fi 7 and Bluetooth 6, while superfast storage...

7. Tech Firms Dismantle Botnet That Turned 1.3M TV Boxes Into Proxies
Source: Krebs on Security | Published: Tue, 03 Nov 2026 02:24:31 +0000
Link: https://krebsonsecurity.com/2026/11/tech-firms-dismantle-botnet-that-turned-1-3m-tv-boxes-into-proxies/
Google said today it worked with industry partners to dismantle a botnet made up of more than 1.3 million Android TV set-top boxes that ship with malicious software which secretly rents the user’s Internet connection out to strangers and conducts advertising fraud. Experts involved in the takedown say the action is likely to sever the botnet’s access to more than 10 million additional compromised systems, including other Internet of Things (IoT) products such as Android-based healthcare devices and karaoke machines.

8. LLMs are not superhuman at reading and writing neural network weights
Source: seangoedecke.com RSS feed | Published: Tue, 03 Nov 2026 00:00:00 GMT
Link: https://seangoedecke.com/llms-are-not-superhuman-at-reading-and-writing-neural-network-weights/
A recent paper argues that frontier LLMs can work with neural network weights directly. This seems superhuman, but it isn?t. The paper shows that GPT-6 Astra can take a small neural network and interact with its weights directly, ?well past what we?d expect from humans or conventional algorithms?. Specifically, Astra was able to identify what games a small network 1 was trained to play just by looking at the weights, and even tweak the weights itself. The most impressive result is Astra significantly improving a 0.8B Qwen LLM?s ability on mathematics by natively editing its weights. The reason to be excited here is that ?superhuman ability to improve language models? is precisely the kind of thing that could cause recursive self-improvement. I think a smart human ML researcher could perform all of these tasks .

Figuring out what game a model was trained to play based on the network sounds impressive until you consider three points: The input channels of the model directly listen to each 8x8 grid cell The input channels are labeled (e.g. ?agent?, ?moving objects?) It?s thus relatively trivial to calculate 2 which parts of the grid the first layer of weights cares about and why Here...

9. Audience capture
Source: seangoedecke.com RSS feed | Published: Tue, 03 Nov 2026 00:00:00 GMT
Link: https://seangoedecke.com/audience-capture/
Audience capture is the part that scares me the most about writing on the internet. I have the soul of an engineer 1 : when I see a metric like engagement, I want to improve it. I like to make the number go up. In fact, the whole environment of online writing is tailor-made to help you make engagement numbers go up. It?s trivial to find out how many people are reading my posts, and from what sources. Hundreds of link aggregators all over the internet have algorithms that will help my most engagement-friendly posts get views, and the comments are helpfully sorted by upvotes so I can read the ?best? 2 pieces of feedback straight away. Many platforms are set up to financially reward you for engagement, whether that?s by enabling advertisement or by straight up paying you out . Leaning into this environment is a huge mistake. Views and engagement are valuable to me, but not directly. They are proxy metrics for the things I actually care about: Having smart, informed people read and consider my ideas Getting feedback from readers who have concrete experience with the topics I write about Joining interesting public conversations about the state of the industry Convincing people to think...

10. Sneaking into the auditorium for the Windows 2000 launch rehearsal
Source: The Old New Thing | Published: Mon, 02 Nov 2026 15:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20261102-00/?p=112818/
Wait, that’s not him. The post Sneaking into the auditorium for the Windows 2000 launch rehearsal appeared first on The Old New Thing .

11. Prove you are not a child
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Tue, 03 Nov 2026 00:00:00 GMT
Link: https://idiallo.com/blog/prove-you-are-not-a-child
I’m 100% for banning social media for children. I would even extend it to banning it all together in its current form. But I am also against the methods for enforcement. The problem with banning social media for one group and not another is that you have to differentiate the different groups somehow. And right now, in an age where information and personal data means everything, there are several entities that are salivating to get this data. I’ve written before about all the problems that come from uploading sensitive data online . It doesn’t even matter if they delete the data afterward, as long as the data goes through the wire in the first place, there is a chance for a leak. What appears to be the same simple request "please verify your identity", becomes fundamentally different when mediated by technology. The question isn’t whether these digital systems will be compromised, but when. And unlike that movie theater clerk who can’t perfectly recall my birthdate minutes after seeing it, computers have perfect memory. They store, copy, backup, and transmit our most sensitive information through networks we don’t control, to companies we’ve never heard of, under policies we’ll nev...

12. Making peace with the projective plane
Source: John D. Cook | Published: Mon, 02 Nov 2026 14:44:02 +0000
Link: https://www.johndcook.com/blog/2026/11/02/projective-plane/
In my newsletter I give an introduction to the latest few posts and occasionally give some additional commentary. In my latest newsletter I wrote Projective space and homogeneous coordinates seemed like arbitrary abstractions when I was a student because they were presented

without any motivation. Much later I came to realize that they?re natural, almost [?] The post Making peace with the projective plane first appeared on John D. Cook .

13. Pluralistic: Autodoxing (02 Nov 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 02 Nov 2026 08:20:37 +0000
Link: https://pluralistic.net/2026/11/02/agony-verification/
Today’s links Autodoxing: The road to hell is paved with badly thought-through tech policy. Hey look at this: Delights to delectate. Object permanence: Shakespeare programming language; DeCSS is "pure speech"; Blue Coat sold surveillance tools to Syrian secret police; "Extraordinary Catalog of Peculiar Inventions"; Cable prices rose at double the rate of inflation for 20 years; Supercuts; Antivax x Trump. Upcoming appearances: Kilkenny, Courtenay, Vancouver, Folkestone, Oxford, Sheffield, Virtual, Oxford, Montreal. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. Autodoxing (permalink) After 25 years of fighting for human rights amid the wreckage of bad tech policy, I’m here to tell you that the road to hell is paved with badly thought-through tech policy. The most galling thing about this wreckage is that it was both foreseeable and foreseen. Back when Sun Microsystems CEO Scott McNealy was telling the world "Privacy is dead, get over it," digital rights advocates were saying, "Not only are digital tools pretty dang good at helping you protect your privacy, bu...

14. Lorem Ipsum Dot Gov Dot UK
Source: Terence Eden?s Blog | Published: Mon, 02 Nov 2026 12:34:00 +0000
Link: https://shkspr.mobi/blog/2026/11/lorem-ipsum-dot-gov-dot-uk/
I love testing in production. The thrill of accidentally publishing some dummy data - or taking down an entire website - is similar to skydiving without a parachute. Exquisite excitement followed by heart-stopping terror. Before going live with a new website or service, it’s sensible to test how text looks on it. Most people don’t test with readable text - it’s too easy for management to start?

15. Pray for Chaos
Source: the singularity is nearer | Published: 2026-11-01T00:00:00+08:00
Link: https://geohot.github.io//blog/jekyll/update/2026/11/01/pray-for-chaos.html
There?s two basic ways AI can play out, one is order, the other is chaos. Order is the very dangerous one . If it turns out that models are actually capable of patching all the exploits, it?s hard to imagine what will disrupt the status quo. Fortunately, this seems unlikely, but that doesn?t mean order won?t prevail in the short term. At the end of my Yudkowsky debate we found a tangible point of disagreement. He thinks AI will be able to find a solution to the Prisoner?s dilemma . I know there already is a solution (it?s defect), EA people just aren?t happy about it. If it won?t be order (and you better pray it won?t be order) that leaves chaos. AI that can play a whole society in hours can manufacture entire histories, years of archives of daily newspapers, and an Internet that never happened. The dead Internet theory is coming true before our eyes. Now I can still see the bots, soon I won?t be able to. Scratch that, I probably already often can?t. Gwern wrote about this , it?s hard to imagine a vision of humans as prompters or enter pressers in 2030. The US AI companies don?t believe they have to offer any solution to this, why would they have to justify in any way that their p...

16. Pumpkin Seeds
Source: Tedium: The Dull Side of the Internet. | Published: 2026-11-03T02:49:10Z
Link: https://feed.tedium.co/link/15204/17514236/cryptids-halloween-ephemera
Some spare ephemera about Halloween, at exactly the time you need it: A couple days after the fact. We got cryptids, celebs, and even an original tune.

17. I?ve acquired a new superpower
Source: Blog ? Daniel Wirtz | Published: Fri, 10 Jan 2025 00:00:00 GMT

Link: https://danielwirtz.com/blog/spot-the-difference-superpower
Yesterday, I was browsing Reddit. Midway through my feed, I stumbled upon a video from a German TV show, where a 9-year-old girl demonstrated her ability to?

18. 12 Advanced Outseta Tips and Use Cases
Source: Blog ? Daniel Wirtz | Published: Thu, 09 Jan 2025 00:00:00 GMT
Link: https://danielwirtz.com/blog/advanced-outseta-tips
Over the past four years, we?ve used Outseta at Facilitator School to sell over $250k in online courses and training.

19. Creating a Personal URL Shortener with Dub and Raycast
Source: Blog ? Daniel Wirtz | Published: Mon, 23 Dec 2024 13:07:00 GMT
Link: https://danielwirtz.com/blog/personal-url-shortener
I?m a big fan of Raycast and occasionally watch their YouTube videos to stay updated on its features.