RSS Digest
Generated at: 2026-11-09 15:42:42

1. The Web Is More Than Enough, by Paul Ford
Source: Simon Willison’s Weblog | Published: 2026-11-09T04:32:33+00:00
Link: https://simonwillison.net/2026/Nov/9/hn-50421091/
My comment on The Web Is More Than Enough, by Paul Ford ? Hacker News. OK, this comment convinced me to give that a shot. Here’s a thing I made in Three.js this morning - a Nudibranch themed Faberge Egg: https://static.simonwillison.net/static/2026/Hidden-Reef-Egg... (77KB) I just asked Claude (Opus 5.5 Medium running in regular Claude Chat) to port that to Vanilla JS, in the same thread where it had built it, and got this: https://static.simonwillison.net/static/2026/Hidden-Reef-Egg... (98KB) They aren’t identical - you can see what Claude said about that in the transcript (I don’t fully understand why it differed): https://claude.ai/share/e56d163e-2793-4677-aa11-191d5058b414 But yeah, it turns out Claude Opus 5.5 is perfectly capable of 3D modeling in Vanilla JS + WebGL without Three.js. In eyeballing the code I find the Three.js slightly easier to understand myself, but I rarely look at the code for these 3D projects so that isn’t a particularly big deal. Tags: hacker-news , javascript , ai , generative-ai , llms , claude , three-js

2. Hide your slop PR descriptions behind
Source: Simon Willison’s Weblog | Published: 2026-11-09T04:25:05+00:00
Link: https://simonwillison.net/2026/Nov/9/details-summary/
Most of the changes I make to my projects these days come in the form of PRs opened for me by Claude Code or Codex. The PR is my unit of work - I use them to bundle commits together, iterating on them until a feature is complete, document decisions and then merge the finished result. Today’s coding agents are great at writing PR descriptions. They’re too good at describing the work The problem is that they are too good. An agent-written PR description is accurate and comprehensive. It describes the problem, the detailed solution, and the steps taken to verify it. That’s a lot of information to absorb. It’s not necessarily slop, but it reads like slop! It’s very easy to spot an AI-written PR, and as with other forms of AI writing my instant reaction is best summarized as "I’m not reading all of that". It doesn’t matter how accurate or useful the information is, I’m instantly put off by the thought that it will take me longer to read than it took the driving author to write... even when that driving author is me! I want human editorial judgement. I want someone to have thought just long enough about the PR to give me a human-approved introduction that helps me understand, at a high...

3. parse_number() as a SQLite extension: C vs Rust
Source: Simon Willison’s Weblog | Published: 2026-11-08T22:26:00+00:00
Link: https://simonwillison.net/2026/Nov/8/sqlite-parse-number-c-vs-rust/
Research: parse_number() as a SQLite extension: C vs Rust I was out walking the dog when I got curious about how easy/difficult it would be to write a SQLite C extension that adds a parse_number() function for turning strings like -$1,234.56 or ’12.5%’ into numeric values that could be summed and sorted. I fired up regular Claude chat (Opus 5.5) and prompted : I’m thinking about a SQLite custom function ’parse_number(value)’ that strips ’,’, ’$’, ’£’ and ’%’, handles a leading ’-’, and returns ’NULL’ for anything it can?t parse Explore the difference between building this as a SQLite extension in C v.s. in Rust - I want to compare performance and code complexity and size of the resulting .so (or whatever) file This gave me a solid starting point - the Rust extension was 321 KB (using Rusqlite) and 298 KB (using sqlite-loadable-rs), while the C extension was 14 KB. But the Rust extension was faster! Turns out that’s because the Rust equivalent of C’s strtod is "twice as fast". So I prompted: Try writing your own C function for parsing text into a decimal number that catches up to or beats the Rust implementation And it did, using "Clinger’s fast path": Net cost per row over 2 mil...

4. Jensen Huang Betrayed Nvidia, Not Just His Family
Source: Daring Fireball | Published: 2026-11-08T23:56:32Z

Link: https://www.nytimes.com/2026/11/06/us/politics/chinese-spy-agency-nvidia-jensen-huang.html?unlocked_article_code=1.O1E.DRyO.9mNcxhaxhtlI
Julian E. Barnes, Tripp Mickle, Paul Mozur, and Maggie Haberman, reporting for The New York Times on Friday: John Rizzo, an Nvidia spokesman, said that the company does not share classified information with its customers, and that customers do not ask Nvidia to press the U.S. government to change its policies on chip sales to China. He added that the Chinese government had ?consistently promoted its own domestic alternatives in China, including Huawei, and strictly limited imports of our data center products.? Mr. Rizzo has maintained that Mr. Huang?s interactions with Ms. Huang were ?not unusual? and consistent with the kinds of meetings he has with any of Nvidia?s many customers. Turns out this was not true, unless Jensen Huang is consistently banging many of Nvidia?s customers. And it now looks especially bad that an Nvidia spokesman was effectively vouching for the Chinese government?s integrity when it comes to their interest in circumventing the ban against selling them high-end Nvidia hardware. Liu Chang, a spokesman for the Chinese embassy, said The Times reporting was ?unsubstantiated? and ?nothing but smears against China.? Turns out this was completely substantiated and...

5. Nvidia CEO Jensen Huang Was Banging a Chinese Spy, Who Was Feeding Information Back to Beijing
Source: Daring Fireball | Published: 2026-11-08T21:50:31Z
Link: https://www.nytimes.com/2026/11/08/technology/nvidia-investigation-jensen-huang-china.html?unlocked_article_code=1.O1E.Soow.6UjrM7Ch-C8L
Tripp Mickle, reporting for The New York Times: Nvidia said on Sunday that it was opening an investigation into Jensen Huang, its chief executive and co-founder, after he acknowledged to the company?s board that he had been intimate with a business executive who U.S. intelligence determined to be a Chinese spy . Mr. Huang?s relationship with the woman, who is named Alice Huang, became a U.S. government concern after surveillance revealed that she was feeding information directly to Beijing. The revelation raised troubling possibilities because Mr. Huang is one of the most powerful figures in the A.I. revolution as well as an influential ally of President Trump. In a message to the company, Mr. Huang apologized. ?I was unfaithful to my wife,? he said. ?My actions hurt her and our children. There is no excuse for what I did. I am deeply sorry.? Nvidia?s board said that it had ?no reason to believe that either national security or Nvidia secrets were compromised,? but that ?the board?s inquiry will examine all available information.? It said Mr. Huang ?fully supports the review and has committed to cooperating with it.? When the NYT first broke this story a little over a week ago, th...

6. Apple?s Analyst Calls Have Been Scripted for Decades
Source: Daring Fireball | Published: 2026-11-03T14:39:03Z
Link: https://9to5mac.com/2026/11/03/apple-ceo-john-ternus-criticized-for-scripted-approach-to-first-earnings-call/
Bloomberg opinion columnist Dave Lee : Apple investors and fans hoping to get an insight into the philosophy of new Chief Executive Officer John Ternus on Monday were left sorely disappointed. The 51-year-old?s first earnings call was a timid affair, with Ternus taking cover behind a carefully written script as analysts peppered him with questions. This is stupid. I mean, it would have been fun if Ternus?s style on the analyst conference call had been loose and freewheeling. Just like it would have been fun if the questions from analysts hadn?t been mostly insipid and pointless. But nothing has changed here other than who the CEO is. Tim Cook?s quarter calls were snoozefests. Save those recordings and use them to help you sleep. And that?s exactly how Cook wanted them. What Ternus should do is follow Steve Jobs?s lead, and skip these calls except when there?s something exceptional to address. Let CFO Kevan Parekh and COO Sabih Khan handle them, the way Jobs used to let Tim Cook and Apple?s then CFOs handle them. At one point in 2008 , Jobs had only appeared on three of the calls in 11 years. He then appeared on the call in October 2008 to address the global financial crisis, and a...

7. LLMs are terrible at designing AI systems
Source: seangoedecke.com RSS feed | Published: Sun, 08 Nov 2026 00:00:00 GMT

Link: https://seangoedecke.com/llms-are-terrible-at-designing-ai-systems/
LLMs don?t know how to design software that uses LLMs. They?re comfortable writing normal programs, but if you throw an AI-shaped problem at them, they will usually screw it up. Why is that? Going down the symbolic rabbit hole For example, this post describes an attempt to vibe-code a memory tool that would flag when one memory contradicted another 1 . This is a great instance of a task that was virtually impossible five years ago but is now relatively trivial due to LLMs. You can simply task a LLM with scanning for contradictions, either in a big batch or pairwise 2 . However, when asked for the ?state of the art? approach 3 , the AI agent in this case did not recommend LLMs. It started building a bewilderingly complex symbolic system to convert the memory items into structured logical atoms , then mechanically compare them with a solver. This, of course, did not and could not work, because there is no analytic way of determining whether any two natural-language sentences express the same proposition. The LLM knew this was impossible! When the user eventually lost patience and said ?hey, why isn?t this working?, Claude immediately explained that this was a known 130-year-old prob...

8. Reverse engineering a vintage diode-transistor logic chip
Source: Ken Shirriff’s blog | Published: 2026-11-08T08:12:25.557-08:00
Link: http://www.righto.com/2026/11/signetics-st680a-dtl-nand.html
Going through a collection of old parts, CuriousMarc encountered some ICs that he couldn’t identify. He opened up one of these ICs, and I studied it under the microscope. After some reverse engineering, I determined that it was a logic chip that implemented four NAND gates using old-fashioned diode-transistor logic (DTL). 1 The chip turned out to be the ST680A, a member of the Signetics 600 series of chips from 1965. In this article, I discuss the circuitry in this chip and how diode-transistor logic works. The die of the ST680A under the microscope. Integrated circuits can be packaged in many ways. One of the most popular is the plastic DIP (Dual In-line Package), a black rectangle of epoxy with two rows of metal pins. This type of package was invented by Signetics: "Introduced in late 1965, the 600 series was the first integrated circuit family to be offered in a solid, plastic, dual in-line package (DIP)."
Earlier integrated circuits were packaged in round metal cans or in small flat-packs. In 1965, engineers at Fairchild invented the DIP, which had two rows of pins on 0.1" spacing. Unlike the surface-mounted flat-pack, the DIP had pins that were inserted into holes, making man...

9. Finite Projective Planes
Source: John D. Cook | Published: Sun, 08 Nov 2026 23:23:49 +0000
Link: https://www.johndcook.com/blog/2026/11/08/finite-geometries/
The idea of finite geometries sounds like an oxymoron the first time you hear of it, something like dry water or cold fire, but the axioms of some kinds of geometry can indeed be satisfied by finite models. The axioms of projective plane geometry are as follows. Any two distinct points lie on exactly one [?] The post Finite Projective Planes first appeared on John D. Cook .

10. Horible public key tricks:
Source: Maurycy’s Blog | Published: Sun, 08 Nov 2026 00:00:00 +0000
Link: https://maurycyz.com/projects/badcert/
So, I’ve always been somewhat annoyed at SSL/TLS. It requires everyone to go through a complex "ACME" dance just to put up a public website. Worse, certificates expire every few months, so it’s not a one and done thing: even with automation, maintaining a certificate requires care to ensure the ACME client always has access to DNS records or the web server. ... and the whole system doesn’t even work because certificates are issued over plaintext connections, and trusted certificate aurthorities use very secure keys like "4" There is prior art of protest SSL implementations , but I do actually want to encrypt web traffic. Passive eavesdropping is a real risk... but preventing it doesn’t require the whole certificate thing: therefore, the best choice is a protest certificate. The only part of the certificate I directly control is the public key. Therefore, the simplest way to obtain a bad certificate is use a vanity private key such as 313333333337 (one of the most eleet primes). However, while the resulting public key

would be week, no one would know unless they actually try to break it. A good protest should be obvious: # openssl s_client -connect maurycyz.com : 443 /dev/null \ |...

11. Some thoughts on restaurant menu QR codes
Source: Terence Eden?s Blog | Published: Sun, 08 Nov 2026 12:34:01 +0000
Link: https://shkspr.mobi/blog/2026/11/some-thoughts-on-restaurant-menu-qr-codes/
Is anyone else SICK AND TIRED of the way hotels now foist plastic "keycards" on guests? No! I don’t want tap to open my door! No! I don’t want to keep the card in my wallet! No! I don’t want electronics in control of my life! What if I need to stab an attacker? Keys are great for that by flimsy plastic cards are useless Gives us PROPER keys with big keyrings that we have to hand in to?

12. Book Review: When the Dust Settles - Stories of Love, Loss and Hope from an Expert in Disaster by Lucy Easthope ?????
Source: Terence Eden?s Blog | Published: Sat, 07 Nov 2026 12:34:34 +0000
Link: https://shkspr.mobi/blog/2026/11/book-review-when-the-dust-settles-stories-of-love-loss-and-hope-from-an-expert-in-disaster-by-lucy-easthope/
This is a difficult but essential book. I was initially put off by how graphically autobiographical it is. The disaster planning I came for was frequently interrupted with the author’s loves, disasters, health issues, pains, and domestic issues. As the book wore on, I realised how vital that is. Disasters - whether natural or caused by human actions - cannot be tackled in a dispassionate way.?

13. Addition By Subtraction
Source: Jim Nielsen?s Blog | Published: Sun, 08 Nov 2026 19:00:00 GMT
Link: https://blog.jim-nielsen.com/2026/addition-by-subtraction/
I still think saying no is important. In that vein, here?s something I shipped by removing functionality (and I think my site is better for it). For the longest time on my icon gallery sites , you could view a list of icons in four pixel dimensions: 56×56 128×128 256×256 512×512 Each of these sizes could be accessed via a control on the website. This has been true for years on my site, whether it?s the current site: Or an older version of the site: By default, I loaded the 128-pixel version of the icons. But if you wanted to see other sizes, you could. However, the more I used the site myself the more I thought, ?Four options is too many?. When it comes to choice, what I really wanted was: Default Smaller Bigger (Plus three is always a good number , unless you have a reason otherwise). So I removed the big 512 option. If I want to study an icon at that size and detail, I click through to the icon?s detail page. Now there are three options for list views: 56, 128, and 256. Small (great on mobile), 128 (great by default), and big (let me see more detail). It?s simple to think about: You get a size by default. Not big enough? Go up a size. Too big? Go down a size. You don?t need to t...

14. Yak-Shaving Overload
Source: Tedium: The Dull Side of the Internet. | Published: 2026-11-09T02:50:32Z
Link: https://feed.tedium.co/link/15204/17518143/yak-shaving-ai-mental-health-effects
Over the years I?ve gotten myself into all varieties of technical rabbit holes and wild goose chases. But I?m feeling weird about my latest trip, involving a forgotten Linux tablet. I blame AI.

15. The great software IP writedown
Source: Martin Alderson | Published: Sun, 08 Nov 2026 00:00:00 GMT
Link: https://martinalderson.com/posts/the-great-software-ip-writedown/?utm_source=rss&utm_medium=rss&utm_campaign=feed
Software’s outsized margins were built on a scarcity of people who could build it. AI agents are removing that scarcity - so where does the value of software IP go?