RSS Digest
Generated at: 2026-11-11 15:33:12

1. llm-gemini 0.36
Source: Simon Willison’s Weblog | Published: 2026-11-10T20:09:44+00:00
Link: https://simonwillison.net/2026/Nov/10/llm-gemini/
Release: llm-gemini 0.36 Gemini 4 Carbon was announced two weeks ago for "a subset of cyber defenders", but today it’s available to everyone else . This is Google’s best ever Gemini, priced at $2/million input and $10/million output at launch (that’s GPT-6.1 Sol / Claude Sonnet 5.5 prices) but scheduled to double in price to $4/$20 on December 1st (same as Opus 5.5). Here are three pelicans for the low, medium and high thinking efforts (it defaults to medium). This is the pelican I got for high: Gemini pelicans have all been great for a while now, so I also tested it out using Google’s Antigravity IDE with a prompt I first tried with Claude Fable 5.2 yesterday : Make a three js animation where yonkers stands up as a big mech (cool stand up motion) and stomps around manhattan and queens and Brooklyn, use openstreetmap mapping data Here’s the HTML result (I had to prompt "build that into a standalone HTML file" to get a version with no external dependencies). I then screen-captured that as a video: Update : Huh, just saw this on YouTube (and tweeted by Sundar Pichai): Using Google Antigravity?s /swarm feature, Carbon built an open-world 3D game where you explore Google?s Mountain V...

2. llm-anthropic 0.31
Source: Simon Willison’s Weblog | Published: 2026-11-09T23:58:46+00:00
Link: https://simonwillison.net/2026/Nov/9/llm-anthropic/
Release: llm-anthropic 0.31 Anthropic released Fable 5.2 and Mythos 5.2 today. Fable 5.2 is currently ranked first on Artificial Analysis . Anthropic say it "solved two open problems in the physics of black holes", so clearly their internal benchmarks have become more ambitious! I’ve added the new models to llm-anthropic , so you can now run prompts through Fable 5.2 like this: uv tool install llm llm install llm-anthropic llm keys set anthropic # paste key here llm -m claude-fable-5.2 "Generate an SVG of a pelican riding a bicycle" I ran pelicans at all five reasoning levels - low , medium , high , xhigh , and max . You can see them all here - here’s the best result I got, from xhigh : max failed to return an image at all, running out of tokens after spending 128,000 thinking about them . That’s the same bug that I saw with Opus 5.5 - but it’s new for the Fable series, both Fable 5 and Fable 5.1 returned max effort pelicans just fine. It’s priced at $8/million input and $40/million output, which is a 20% discount on the previous Fable 5.1 ($10/$50). I also used Fable 5.2 to help identify the source of a mysterious beep . Tags: llm-release , llm , anthropic , claude , generative-a...

3. Locating the source of a mysterious beep with Claude
Source: Simon Willison’s Weblog | Published: 2026-11-09T23:57:21+00:00
Link: https://simonwillison.net/2026/Nov/9/mysterious-beeps/
Our house has a box room, full of boxes that we still haven’t unpacked from our last move. A couple of weeks ago something in that room started intermittently beeping . The beeps were far enough apart that standing in the room and trying to determine the source of the noise was frustrating, and I really didn’t want to dig through every one of those boxes. So I had Claude Code build Beep Locator , a software system for triangulating the source of a mysterious beep using multiple devices scattered around the room. Here’s what Beep Locator looks like after it has been running for a while: You run Python web server to show the primary UI, and then use the QR code or URL to connect as many devices as you can spare to act as listeners - I was using two iPads, a phone, a MacBook Pro and a Surface Ultra. The devices then play noises to each other to help estimate their locations. This didn’t work particularly well, so I added a "Rough mode" where you can drag their icons around in the space to provide rough locations which are then adjusted based on those sounds. Once that’s set up, the devices all listen for the beep! I added a clapping mode too so I could test it by clapping my hands in...

4. Microsoft Fixes Nearly 1,100 Security Vulnerabilities

Source: Krebs on Security | Published: Tue, 10 Nov 2026 21:55:03 +0000
Link: https://krebsonsecurity.com/2026/11/microsoft-fixes-nearly-1100-security-vulnerabilities/
Microsoft did another four-digit patch drop today, releasing updates to fix almost 1,100 security vulnerabilities in its Windows operating systems and supported software. This month’s Patch Tuesday addresses 63 critical flaws -- including a zero-day in Microsoft Outlook -- and comes amid important security updates from other software makers, including Adobe, Apple, Google and Mozilla.

5. Gary Marcus: ?Just Say No to Anthropic?s IPO?
Source: Daring Fireball | Published: 2026-11-10T22:16:25Z
Link: https://garymarcus.substack.com/p/just-say-no-to-anthropics-ipo?r=3k69a0
Gary Marcus: When a business needs exponentially increasing amounts of capital to make modest increases in quality in an intensely competitive market and still owes far more money than it has ever made, there is a serious problem. I continue to believe that there is no moat in AI models. None. But the entire premise of Anthropic?s IPO is that they have an enormous gulf of a moat. I would not be shocked to see Anthropic go out of business. Nor would I be shocked if the IPO fails. WeWork tried to IPO, 3 times, and didn?t make it. I would be surprised if Anthropic does not go out of business. Most cults implode, and Anthropic doesn?t even hide that they?re a cult . And I just don?t see how the math works. Scientology is a long-running cult, but Scientology hasn?t committed $900 billion in spending over the next few years. (Marcus?s post is on his Substack blog, which, annoyingly, requires you to be logged into a Substack account and to ?subscribe?, even though it?s ostensibly a free post. What a rotten oppressive turd of a blogging platform Substack is. Just in case your Substack account isn?t working, here?s a PDF archive .) ?

6. Altar II Keyboard on Kickstarter
Source: Daring Fireball | Published: 2026-11-10T20:59:25Z
Link: https://www.kickstarter.com/projects/electronicamaterials/altar-ii/
Electronic Materials Office is the brainchild of Andrew Emerson. Back in 2023 he launched the Altar I, a rather distinctive mechanical keyboard. Much thinner than most, very opinionated in its industrial design and keycap typography. I got one and I don?t regret it, but there was too much I didn?t like about it, so it?s mostly sat in a drawer. Today on Kickstarter Emerson/EMO has launched a campaign for the Altar II. Same cool aesthetic, but with improvements that address everything I didn?t like about the Altar I. The Altar II has all new switches. It?s so thin it resembles the profile of Apple?s Magic Keyboard from the side. The dial ? now called the M-Dial ? now attaches magnetically, which means it?s easy to detach for putting the keyboard into a bag. And it has new features that never even occurred to me, like NiteLite ? backlighting from ?a long-wavelength amber-red light that won?t affect your body?s circadian rhythm, meaning you can sleep better at night.? The sleep benefits might be voodoo but red lights are undeniably easier on your eyes at night, and it looks cool as shit. (My favorite feature of Apple Watch Ultra is ?Night Mode? ? a handful of watch faces that switch t...

7. Bloomberg: ?Nvidia Intervened in U.S. Probe of Megaspeed Chip Smuggling?
Source: Daring Fireball | Published: 2026-11-10T17:06:01Z
Link: https://www.bloomberg.com/news/articles/2026-11-09/nvidia-intervened-in-us-probe-of-megaspeed-chip-smuggling
Mackenzie Hawkins, reporting yesterday for Bloomberg: Despite that stance, in the months after Teter?s initial call with Kessler, Nvidia gave BIS an incomplete account of its business with Megaspeed. After several visits to Megaspeed?s data centers ? 12 facilities owned by eight companies across three countries in Southeast Asia ? Nvidia had sent inventory checks to the agency and declared to Bloomberg News in mid-November 2025 that everything was where it should be. The company also said at the time that it had asked the US government whether it needed to cut off business with Megaspeed, and was told no. But Nvidia?s tabulations, as relayed to the US government, undercounted the number of chips it had sold to Megaspeed via hardware partners by 50,000 units, according to trade records analyzed by Bloomberg. That?s worth billions of

dollars. After being pressed about the gap, Nvidia told Bloomberg in December 2025 that it had checked additional Megaspeed warehouses that contained the missing hardware, while declining to state exactly how many chips it found in those facilities or provide any further details. ?

8. Glue work is still harmful in the age of AI
Source: seangoedecke.com RSS feed | Published: Wed, 11 Nov 2026 00:00:00 GMT
Link: https://seangoedecke.com/glue-work-is-still-harmful-in-the-age-of-ai/
I am not a believer in glue work 1 . Processes that can only be sustained by the unrewarded sacrifice of engineers ought to die. Now that LLMs are writing most of our code, does that free human engineers up to do more glue work ? No, it doesn?t. I think glue work is actually less important to companies than it was in the 2010s. Let?s go through four reasons why. First, AI decreases the amount of coordination needed to ship . The magnifying effect of AI agents means that any single engineer can now take on projects that would previously have needed a small team. A lot of glue work was helping with coordination: booking kickoff and status meetings, passing around internal updates, and so on. A single engineer doesn?t need to do any of that. Second, engineers in the 2020s require less ?careful handling? . In my experience, the most important subcategory of coordination-based glue work was managing difficult people . Many companies have secret gatekeepers: engineers with outsized influence and large egos who will shoot down your project if they feel insufficiently consulted. I?ve written before about how some teams are structurally set up to be gatekeepers, and how some engineers coul...

9. Microsoft employee price list from 1987
Source: The Old New Thing | Published: Tue, 10 Nov 2026 15:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20261110-00/?p=112843/
The Microsoft product catalog. The post Microsoft employee price list from 1987 appeared first on The Old New Thing .

10. Balanced Incomplete Block Designs
Source: John D. Cook | Published: Tue, 10 Nov 2026 14:26:21 +0000
Link: https://www.johndcook.com/blog/2026/11/10/bibd/
Suppose you have n different fertilizers you want to test. The simplest experiment would be to take a big plot of land, divide it into n sections, and apply each fertilizer to a section. That?s fine if you have one big enough plot of land, and the land itself doesn?t matter. If the fertilizer supplies [?] The post Balanced Incomplete Block Designs first appeared on John D. Cook .

11. Finite Affine Planes
Source: John D. Cook | Published: Mon, 09 Nov 2026 14:21:49 +0000
Link: https://www.johndcook.com/blog/2026/11/09/finite-affine-planes/
The previous post was on finite projective planes. This post will look at finite affine planes and the next will look at finite hyperbolic planes. Affine geometry is essentially what remains of Euclidean geometry if you keep the incidence properties, e.g. wanting to know whether things intersect, and forget about metric properties such as segment [?] The post Finite Affine Planes first appeared on John D. Cook .

12. Book Review: The Unfinished Harauld Hughes - Richard Ayoade ?????
Source: Terence Eden?s Blog | Published: Tue, 10 Nov 2026 12:34:31 +0000
Link:
https://shkspr.mobi/blog/2026/11/book-review-the-unfinished-harauld-hughes-richard-ayoade/
A decade ago, I read Ayoade’s first book. In my review, I said: It’s hard to know how to assess a book like this. It’s far too clever for its own good - a frustrating and humorous look at... what? It’s neither autobiography nor comic novel. An anti-novel which trudges the reader through a series of comic vignettes. It reminded me of some of Woody Allen’s work. If you get the joke - you’ll find ?

13. Pluralistic: The EU should go big and hard over border IT failure (10 Nov 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 10 Nov 2026 17:34:21 +0000
Link: https://pluralistic.net/2026/11/10/worst-of-all-worlds/
Today’s links The EU should go big and hard over border IT failure: Sometimes you have to execute an IT contractor to encourage the others. Hey look at this: Delights to delectate. Object permanence: Toothpick sculptures; McRib as a commodities volatility hedge; Hitler’s British passport; Why we get tech wrong; France jails spying EDF execs; SF v Trump; Clinton’s data-driven ground game sucked; Risk compensation sucks; "Flexible labor" is risk compensation. Upcoming appearances: Vancouver, Folkestone, Oxford, Sheffield, Virtual, Oxford, Santa Barbara, Montreal. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. The EU should go big and hard over border IT failure (permalink) If the EU has a future as a political project, it will have to be as a living, breathing repudiation of Ronald Reagan’s maxim that "The nine most terrifying words are ’I?m from the government and I?m here to help.’" In other words, the EU’s future depends on its ability to publicly perform competence and capacity. The EU is a technocratic project, and the one trait technocrats must show is un...

14. 20 years of blogging
Source: Giles’ blog | Published: Tue, 10 Nov 2026 23:00:00 +0000
Link: https://www.gilesthomas.com/2026/11/20-years-of-blogging
20 years ago to this day, I started my new technical blog. I’d had a personal one for some time, but most of the stuff that happened there had already migrated to social media. I wanted somewhere to post the things that would only be of interest to my techier friends, and maybe of some use to anyone else that came across it. Things have changed a bit, but they’ve stayed the same in some ways. Earlier this year I finished my 48-part (!) LLM from scratch series, but it had predecessors -- the third post I made here was part one of a 13-part series where I worked through setting up automated backups for my homelab NAS, and then later on I spent a few months building a reverse proxy in C , not to mention the eight-month, ten-post series on fine-tuning LLMs . It does seem like there’s some value to putting all of this online. I rarely check my logs and referrers, but people do still seem to be arriving here having searched for stuff that I’ve written about, and every now and then there’s a kind "thanks for this!" comment left on an old post where I explained something I’d just learned. I wrote a bit more about the blog’s history last year, and came to the conclusion that the best posts...

15. Jensen’s Bizarre Adventure
Source: Ed Zitron’s Where’s Your Ed At | Published: Tue, 10 Nov 2026 17:43:27 GMT
Link: https://www.wheresyoured.at/jensens-bizarre-adventure/
If you liked this piece, you should subscribe to my premium newsletter, and you can subscribe on the following links: $70 a year , $18 a quarter , or $7 a month . In return, you get a weekly premium newsletter that dives deep into the finances of the AI

16. Making a Website in 2026
Source: Jim Nielsen?s Blog | Published: Tue, 10 Nov 2026 19:00:00 GMT
Link: https://blog.jim-nielsen.com/2026/making-websites-in-2026/
Pete Millspaugh has a great post over on the Val Town blog where he documents his attempt at making generating the same website across a variety of AI tools. I love the post. Here are a couple of thoughts that come to mind. The Space Is Ripe for Referral Fees Pete decided to make a website for his conference talk, and he added a comments section to make it a less static / more dynamic demo. This is interesting because LLMs are really good at generating one-off HTML pages you open at file:// , but as soon as you start to introduce technological dependencies ? like a database, or an analytics tool, or a service for sending email, or payment processing, etc. ? the question becomes: How does the LLM choose which service to rely on? It?s one thing to have the LLM choose React over Angular or Vue over Svelte. But how (and why) does it choose something like an auth provider? Does it build you a site with Auth0 or WorkOS or Clerk? How

would a non-technical user even know how to weigh the trade-offs of making such a decision? Pete said he was deliberately vague about these details, which is why I find his experience so intriguing. For example, here?s what he found after generating a site i...

17. Claude Code token cost explorer
Source: Entropic Thoughts | Published: Tue, 10 Nov 2026 00:00:00 +0100
Link: https://entropicthoughts.com/claude-code-token-cost-explorer
I am experimenting with involving subagents more in my LLM coding agent workflow. I discovered that subagents make it more difficult to diagnose token usage in Claude Code, because the regular /context command only breaks down context usage for the main conversation, not for the subagents. I tried some existing tools like ccusage and claude-agent-ledger , but they weren?t able to give me the sort of breakdown into specifics I needed to optimise my usage. Thus, I shamelessly had Claude Code whip up a tool that gives me a better view of its token usage. This is vibecoded in the true sense of the word ? I have barely looked at the code ? but I?m okay with that becuase it is not critical infrastructure and it does what I need it to. If you want to know how to get more out of your Claude Code sessions without paying more for it, here?s how you can find out. First, open the context cost explorer in your browser. Everything happens locally on your machine, so none of your transcript data leaves your computer. If you?re the paranoid kind, you can verify by opening the network tab in your browser and reading in a non-sensitive session and seeing that nothing funny happens. Then follow the...

18. Monads in Lean
Source: Eli Bendersky’s website | Published: 2026-11-10T19:25:00-08:00
Link: https://eli.thegreenplace.net/2026/monads-in-lean/
This post demonstrates how to use monads in the Lean programming language (Lean 4, to be specific). If you’re not familiar with Lean, this part of the language is very similar to other functional programming languages, like Haskell or OCaml. Motivating example - evaluating expressions Suppose we want to write a ?

19. Finally, Payment by Bank Transfer Comes to Have I Been Pwned
Source: Troy Hunt | Published: Tue, 10 Nov 2026 20:26:15 GMT
Link: https://www.troyhunt.com/finally-payment-by-bank-transfer-comes-to-have-i-been-pwned/
This is simultaneously the most boring yet exciting blog title I’ve written in a long time. How you pay for services isn’t exactly a sexy topic, but it solves a problem that’s been a massive bugbear for a long time. Here’s a

20. I am searching for a coffee alternative
Source: Max Bernstein’s Blog | Published: Tue, 10 Nov 2026 00:00:00 +0000
Link: https://bernsteinbear.com/blog/coffee-alternative/?utm_source=rss
I had an in-person work team meeting last week. In the office, they had Bevi seltzer machines. When Bevi machines first arrived at Facebook, where I used to work, I looked down on them. They seemed to be a cost reduction, a sad imitation of The Good Stuff that came in cans and bottles in the fridge. Good Stuff such as LaCroix 1 and Ito En tea 2 . But this week was different. I noticed that this Bevi had a caffeine option 3 . After some experimentation, I decided that for the week, I would drink unflavored, chilled, sparkling, caffeinated water instead of coffee. And it was great. Basically none of the coffee side effects. For several days straight I thoroughly enjoyed a 100-120mg caffeine sparkling water. But all good things must come to an end. You cannot (yet?) buy a Bevi machine for home. So I am unfortunately now searching for alternatives. I took a stroll through our local Star Market in search of a replacement. I looked online, too, since we have a weird selection of beverages at that market. I ended up with an assortment of brightly colored cans. Time to do a taste test. Requirements In order to swap out coffee for something else, I have the following criteria: no (gratuito...

21. We need more goat sacrifices
Source: Experimental History | Published: Tue, 10 Nov 2026 17:21:43 GMT

Link: https://www.experimental-history.com/p/we-need-more-goat-sacrifices
OR: OJ A-OK

22. Revisiting Leander v Sweden and Big Brother Watch
Source: Bert Hubert’s writings | Published: Tue, 10 Nov 2026 20:10:00 +0100
Link: https://berthub.eu/articles/posts/revisiting-leander-and-big-brother/
Governments in Europe are gearing up to increase their bulk surveillance capabilities. One check on their ambitions is the European Convention on Human Rights. In this piece, I want to highlight recent jurisprudence from the European Court of Human Rights (ECtHR), and interpret the case law in light of current cybersecurity and AI developments. This interpretation may have relevance for future evaluations of SIGINT and other bulk data collection programs. Because what seemed reasonable a few years ago may no longer be so.