RSS Digest
Generated at: 2026-11-13 15:39:36

1. llm-gemini 0.37
Source: Simon Willison’s Weblog | Published: 2026-11-12T22:55:19+00:00
Link: https://simonwillison.net/2026/Nov/12/llm-gemini/
Release: llm-gemini 0.37 Google launched Gemini 4 Hydrogen today - the lowest cost member of the Gemini 4 family. It’s actually not too far off the price of its much bigger siblings, Carbon and Argon. I’m finding Gemini 4’s pricing really confusing. Here’s the pricing scheme: Model Dates Input ($/M) Cached input Output Gemini 4 Carbon Through Nov 30, 2026 $2.00 $0.10 $10.00 From Dec 1, 2026 $4.00 $0.20 $20.00 Gemini 4 Argon Through Nov 30, 2026 $2.00 $0.10 $10.00 From Dec 1, 2026 $4.00 $0.20 $20.00 Gemini 4 Hydrogen $1.50 $0.075 $7.50 Carbon and Argon are the same price , and both are scheduled to double in price on December 1st. If they’re the same price, why have two different models? Google says that Argon "delivers Gemini 4 intelligence with faster responses and lower token spend per turn compared to Gemini 4 Carbon" - so it may be priced the same, but it will use less tokens per task. Hydrogen, supposedly the cheap model, is currently only a 25% discount on its more capable peers. The Hydrogen price will make more sense after December 1st, when it will represent a 75% discount. It’s still considerably more expensive than Claude Haiku 5.5 and GPT-6 Luna - they’re both $0.10/$0...

2. llm 0.37.2
Source: Simon Willison’s Weblog | Published: 2026-11-11T17:19:52+00:00
Link: https://simonwillison.net/2026/Nov/11/llm/
Release: llm 0.37.2 Fixes a security issue reported by Daniel de Araujo in GHSA-v9h7-jj4g-f862 and GHSA-r639-rr7r-x72w . LLM has a mechanism called extra-openai-models.yaml - now managed by the new-in-0.37 llm openai endpoint add command - for using LLM with OpenAI-compatible API endpoints from other providers. Daniel reported two ways this could accidentally send your configured OpenAI API key to one of those providers: if the key named by api_key_name had never been stored, or if a responses: true model was run with -o chat_completions 1 . Tags: projects , security , llm , releases

3. Anthropic’s IPO roadshow video
Source: Simon Willison’s Weblog | Published: 2026-11-11T15:11:59+00:00
Link: https://simonwillison.net/2026/Nov/11/anthropic-ipo-roadshow/
Anthropic’s IPO roadshow video Anthropic kicked off their IPO roadshow yesterday, and their roadshow video has some previously unreleased information about how their business works. The section on "Approach to compute" with co-founder Tom Brown is particularly interesting, describing how Anthropic train and serve Claude across AWS Trainium, Google TPUs, NVIDIA GPUs, and AMD GPUs. I’m not sure why there’s no official video transcript, but here’s one that I extracted using MacWhisper. Highlights below: We train and serve Claude now across all four major chip families. That’s only possible because early on we made an unconventional bet on a multi-chip strategy. At the time, many people thought that running AI models on different chips wasn’t worth the effort. Each chip needs its own tailored software stack and inference stack, which is expensive and difficult to build. We’ve now been doing it for years, and we’ve gotten incredibly good at it. [...] As an example, this year, we started working with AMD GPUs for the first time. In the past, bringing on a new chip platform took dozens of engineers months of work This time, we had Mythos try it out. It got Claude running on AMD in just t...

4. So you want to buy AI for your company
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Fri, 13 Nov 2026 00:34:41 GMT
Link: https://idiallo.com/byte-size/buying-ai-for-your-company
Where do I start? We have Gemini 4 Hydrogen. But it sounds like you could use Claude Sonnet 5.5. Or maybe 5.4; it really depends. Have you tried Opus? We have Opus 5.2 right now. It’s

similar to Fable 5.2, but better for day-to-day work. It’s not quite as powerful as Mythos 5.2, though. I hear the gap will narrow eventually. Opus scores 94.8 on a benchmark, while Mythos gets 96.1. Fable is a respectable 92.7, assuming you don’t ask it to do anything too ambitious. Have you tried the Chinese models? DeepSeek is pretty good. You might want to start with the R1 family to test the waters before moving on to V4. Personally, I’m a Flash kind of guy, but Pro is pretty good too. The crowd favorite is GLM these days. You’ve got to try 5.3 Flash. Its Mixture-of-Experts architecture is state of the art, in my opinion. It scores an impressive 87.3 on the Benchmark Performance Index. But let’s be serious for a minute. You’re probably looking for ChatGPT. Astra is overkill unless you want to do some damage. I’d stick with GPT-6.0 Terra, not Sol! That one ate through my budget in three days. Terra gets 98.2 on the Enterprise Productivity Benchmark, while Sol gets 99.1. That extra 0.9 will cost y...

5. How can I pass more than one parameter to co_await?
Source: The Old New Thing | Published: Thu, 12 Nov 2026 15:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20261112-00/?p=112851/
The co_await operator takes only one parameter, but it can be a big parameter. The post How can I pass more than one parameter to co_await ? appeared first on The Old New Thing .

6. Jensen Huang Agrees to Testify Before Congress, But Not Publicly (Yet)
Source: Daring Fireball | Published: 2026-11-12T21:32:07Z
Link: https://techcrunch.com/2026/11/12/jensen-huang-agrees-to-testify-before-congress/
Aisha Malik, TechCrunch: Ex-Nvidia CEO Jensen Huang has agreed to testify before Congress, House Foreign Affairs Committee Chairman Brian Mast said on Thursday, according to reports from Bloomberg and NBC News . Huang, who stepped down as CEO following reports of an alleged affair with a woman suspected of spying for China, agreed to testify Wednesday after several senior lawmakers threatened to subpoena him. Huang has agreed to testify before Congress by Thanksgiving, which falls later this month. [...] The testimony will take place behind closed doors, though Mast said he expects the briefing to be recorded on video and transcribed. Closed-door is pantomime. Public reckonings will come next year . This whole thing could crater Nvidia?s stock price, sinking the market, and it?s Nvidia?s investments that are propping up the entire AI industry. So if the whole economy takes a severe beating in the coming months, imagine just how much more contentious Congressional inquiries under Democratic control could be. I?m of the belief that the whole AI boom is ultimately a complex Ponzi scheme, and Ponzi schemes, by their nature, eventually crash out. But now Huang has set himself up as the...

7. Scott Colom Flips Mississippi
Source: Daring Fireball | Published: 2026-11-12T21:20:29Z
Link: https://talkingpointsmemo.com/live-blog/new-congress-gets-underway-with-hegseth-impeachment-effort
Emine Yücel, reporting for TPM: Mississippi prosecutor Scott Colom officially defeated incumbent Sen. Cindy Hyde-Smith (R-MS) Wednesday night. Colom will be the first Democrat to represent Mississippi in the Senate since 1982 and he will also become the state?s first Black senator in nearly 150 years. 150 years is, needless to say, a long time. The crazy part about that is that Mississippi has the highest proportion of black citizens in the entire country (38 percent). Colom, by all accounts, ran a fantastic campaign, with little to no support from the national Democratic Party because almost no one imagined Mississippi ? the deepest of the deep South ? might turn blue. This commercial Colom ran is one of the best I?ve ever seen. Not just best political commercials. I mean it?s one of the best 30-second spots ever made. It?s perfect. I just keep watching it over and over. Update: I stand corrected on the national support. The WSJ ran a story back in September ? six weeks before the election ? on Senate Minority Leader Chuck Schumer raising funds for Colom amongst big Democratic donors. Schumer smelled a winner. ?

8. The Dream of Quibi Lives On
Source: Daring Fireball | Published: 2026-11-12T20:54:43Z

Link: https://www.nbcuniversal.com/article/peacock-extends-its-vertical-video-leadership-reimagined-mobile-experience
Actual headline from an actual press release: ?Peacock Extends Its Vertical Video Leadership With a Reimagined Mobile Experience?. Please tell me that someone printed up some t-shirts and hats with ?Vertical Video Leadership? on them. ?

9. Book Review: This Is for Everyone - Tim Berners-Lee ?????
Source: Terence Eden?s Blog | Published: Thu, 12 Nov 2026 12:34:10 +0000
Link: https://shkspr.mobi/blog/2026/11/book-review-this-is-for-everyone-tim-berners-lee/
Is this a biography of a man, or of his mission? Is this the story of how the web came to be, or a lament for what happened to it? Do you crave the inside gossip, or the PR sanitised gloss? This Is For Everyone, but I think it ends up being slightly unsatisfying for anyone. It starts with the normal charming family anecdotes. No grand revelations or hidden secrets, just a lovely ramble?

10. Technology as a Search Problem
Source: Construction Physics | Published: Thu, 12 Nov 2026 13:01:45 GMT
Link: https://www.construction-physics.com/p/technology-as-a-search-problem
One of my ongoing research interests is understanding the nature of technological progress: How does new technology come into the world and improve over time?

11. First Strike
Source: the singularity is nearer | Published: 2026-11-13T00:00:00+08:00
Link: https://geohot.github.io//blog/jekyll/update/2026/11/13/first-strike.html
Internalize this graph, we post it on this blog a lot. If trends continue, China will have 10x the electricity of the US in 10 years. Whenever you think about geopolitical strategy, always keep it in mind, because whatever your plans are must consider this. Here?s a strategy that would make sense for America long term. Sell as many AI chips into China as possible, use the money to purchase solar panel and battery factories, and build out electricity capacity as quickly as possible. Fundamentally, compute is a function of electricity. That?s the sort of long term thinking that gives America a chance. With robotic labor, immigrants like Jensen and Elon, a business friendly culture, and an incredible chunk of land, it?s not so clear who is ahead in 50 years. Regardless, I do know this strategy would lead to massive prosperity for the citizens of both countries. However, that?s not close to what we see happening. We see a lot of short term zero sum thinking in action. It?s very scary what the US ?AI race? mentality leads to. When exactly does the race end? Whatever capabilities you have today will eventually be matched, unless? ?you believe that you can use ASI in a decapitating first...

12. Pluralistic: Uber’s labor crash (12 Nov 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Thu, 12 Nov 2026 20:50:53 +0000
Link: https://pluralistic.net/2026/11/12/starvation-wages/
Today’s links Uber’s labor crash: Replacing serfs with robots just doesn’t pencil out. Hey look at this: Delights to delectate. Object permanence: Asimov inspired Al Qaeda; PATRIOT Act scares Canadian customers off US servers; EFF v Trump; Internet infrastructure is unsupported and failing; Nominative determinism debunked; How to fight Trump; Burglars v burglary; NYT said Hitler was only kidding; Folky Song Funny Note; Surviving autocracy; Corporate criminals. Upcoming appearances: Vancouver, Folkestone, Oxford, Sheffield, Virtual, Oxford, Santa Barbara, Montreal. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. Uber’s labor crash (permalink) There’s a specific kind of lie that tech bosses tell: the "premature truth." That’s a lie that they excuse by saying it’ll come true soon: https://pluralistic.net/2024/12/04/its-not-a-lie/#its-a-premature-truth Elon Musk is the world’s most enthusiastic teller of these lies. That’s not surprising: any kind of radical promise about the automotive sector has a 90% chance of being this kind of bullshit. Musk can

barely mentio...

13. SSA from bytecode: imperative pearl
Source: Max Bernstein’s Blog | Published: Thu, 12 Nov 2026 00:00:00 +0000
Link: https://bernsteinbear.com/blog/ssa-bytecode/?utm_source=rss
Ufuk made a comment that he enjoyed my description of SSA from bytecode in Aaron?s livestream with me . While filming the livestream?an intro to ZJIT internals?I got a little nervous that people would shy away from our 1900 line add_iseq_to_hir function. This function is responsible for converting CRuby?s stack based bytecode for the YARV virtual machine into a high-level SSA IR for ZJIT to analyze and optimize. It takes bytecode that looks roughly like this: push 3 push 4 add into SSA form, giving an SSA value to each stack slot, ish: v0 = const 3 v1 = const 4 v2 = add v0, v1 While our function is 2000 lines long because it handles a lot of specific opcodes and weird facts about those opcodes, the core of it is small. Here is that core. What if a stack, but virtual The bytecode assumes that when it is run, it will be operating on a stack. This means that every opcode has implicit stack inputs and implicit stack outputs. push 3 , for example, has empty stack inputs and one stack output. By contrast, add has two stack inputs and one stack output. Some instructions such as nop are empty on both counts. As far as I can tell, all instructions have at most one output. These counts are...